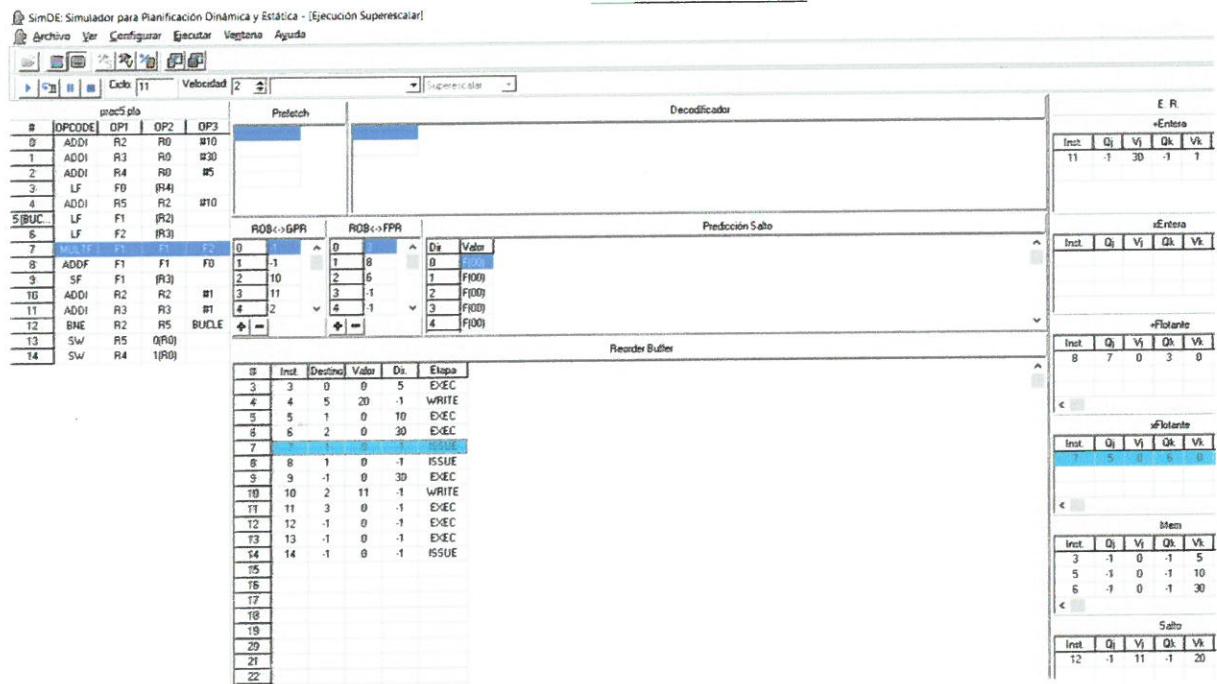


## Análisis de las columnas en la ER:



### 1. Columna $Q_j$ (Etiqueta de Dependencia del Operando 1):

- El valor de  $Q_j = 5$  indica que el operando  $j$  aún está pendiente de ser producido por una instrucción previa (la instrucción 5 en este caso). Esto implica que no está listo para ejecutar porque depende de los datos generados por esa instrucción.  
*No dep*

### 2. Columna $Q_k$ (Etiqueta de Dependencia del Operando 2):

- El valor de  $Q_k = 0$  indica que el operando  $k$  no tiene dependencias y está listo para usarse.

### 3. Columna $V_k$ (Valor del Operando 2):

- El valor  $V_k = 6$  confirma que el operando  $k$  tiene un valor disponible, lo cual corrobora que no hay problemas con este operando.

### 4. Columna $V_j$ (Valor del Operando 1):

- El valor  $V_j = 0$  indica que el operando  $j$  no tiene un valor disponible porque depende de la instrucción identificada en  $Q_j$  (la instrucción 5).

### 5. Columna $A$ :

- Aunque no es directamente relevante en este caso, el valor  $A = -1$  puede indicar una dirección de memoria o desplazamiento pendiente, lo cual no afecta la ejecución actual de esta instrucción.



## LAB ASSIGNMENT 3

### STATIC CODE SCHEDULING

#### Objectives

- Understand the *loop unrolling* technique.
- Understand the *software pipelining* technique.

#### Procedure

The student must apply different static code scheduling techniques and evaluate by simulation the performance of the resulting code, assuming pipelined multicore processors.

Consider the SAPY or DAPY operation (Single or Double precision A Plus Y), consisting in computing  $Z = a + Y$ , where Y and Z are vectors stored in memory, and a is a scalar value.

A possible implementation for the *SimulazMS* simulator is given in Appendix A. Configure the simulator (version 4.12i) with 4 floating-point addition units with a latency of 4 cycles and leave the default configuration for the rest of functional units.

Then, do the following tasks and answer to the corresponding questions:

- For the SAPY program:
  - Check its behavior.
  - Compute the CPI (clock cycles per instruction) and CPR (cycles per result) values.
- Apply loop unrolling. Unroll the minimum number of times to eliminate any stalls.
  - Check the behavior of the modified code.
  - Compute the CPI and CPR values.
  - Compute the obtained speedup (with respect the original code).
- Apply software pipelining.
  - Check the behavior of the modified code.
  - Compute the CPI and CPR values.
  - Compute the obtained speedup (with respect the original code).

#### APPENDIX A

```
# SAPY → Z = a + y
```

```
# data segment
```

```
.data
```

```
# Vector size = 64 words = 256 bytes
```

```
y:
    .float 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18,
    19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38,
    39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58,
    59, 60, 61, 62, 63
```

```
# Vector z
```

```
.space 256
```

```
a:
```

```
# Scalar a
```

```
.float 1
```

```
# text segment
```

```
.text
```

```
.globl main
```

```
main:
```

```
la $t1, y
    # $t1 contains the address of y
la $t2, z
    # $t2 contains the address of z
la $t3, a
    # $t3 contains the address of a
```

```
lwcl $f0, 0($t3)
addi $t4, $t1, 256
    # f0 contains a
```

```
lwcl $f2, 0($t1)
addi $t4, $f0, $t2
    # $t4 contains the address of f0
swc1 $f4, 0($t2)
addi $t1, $t1, 4
addi $t2, $t2, 4
bne $t4, $t1, loop
addi $v0, $0, 10
    # end
syscall
```

```
58S
461
$ no depend.
$ no depend.
```

*Loop unrolling:*

```
lwcl f2, 0(f1)
lwcl f12, 4(f1)
add f4, f0, f2
add f14, f0, f12
add t2, t2, 8
add t4, t4, 8
swcf q1, -8(t2)
```

1

1. *No syscall* → CPI =  $\frac{\text{Ciclos}}{\text{Instrucciones}} = \frac{780}{8 + (64 - 7)} = 1'710S$

CPR =  $\frac{\text{Ciclos}}{\text{Vuelta}} = \frac{780}{64} = 12'187S$

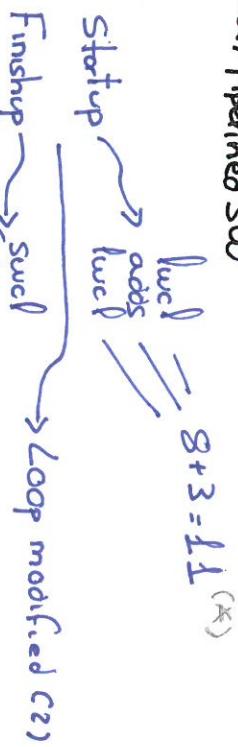
Speedup =  $\frac{\text{Ciclos}}{\text{Ciclos Nuevo}} = \frac{780}{64} = 12$

Speedup =  $\frac{\text{Ciclos}}{\text{Ciclos Nuevo}} = \frac{780}{64/2} = 21591$

### 3. Segmentación / Pipelined SW

Modified loop → No stalls

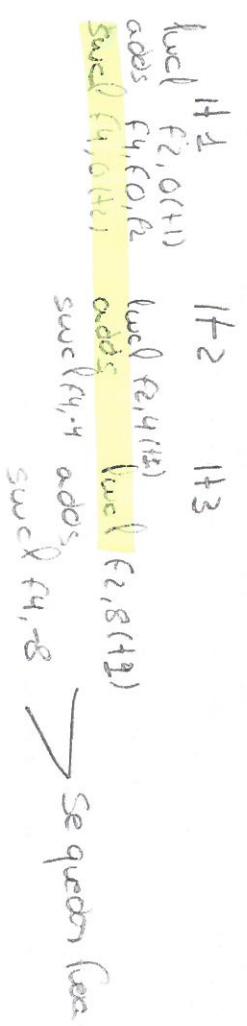
Modifico loop → solo rebocado → problemas??



$$\text{CPI} = \frac{397}{12 + (62 \cdot 6) + 4} = 1.05685$$

$$\text{CPR} = \frac{397}{64} = 6.13906$$

$$\text{Speed-up} = \frac{780}{397} = 1.971$$



**QUESTIONS:**

**1. History and Context**

**Q: What is the historical significance of the Intel 8086 microprocessor?**

**A:** The Intel 8086, released in 1978, is significant for being the foundation of the x86 series of processors, which have become a standard in personal computers. It paved the way for modern microprocessor designs.

**Q: How does the Meteor Lake architecture represent a shift for Intel?**

**A:** Meteor Lake introduces a chiplet-based design, moving away from the traditional monolithic design. This allows for better modularity, scalability, and integration of AI and other capabilities.

**2. General Characteristics**

**Q: What is the purpose of Performance (P) and Efficiency (E) cores in the hybrid architecture?**

**A:** P-cores handle demanding, intensive tasks such as 3D rendering and gaming, while E-cores optimize background and low-priority tasks, improving overall efficiency.

**Q: How does the Foveros 3D stacking technology contribute to processor performance?**

**A:** Foveros 3D allows for stacking chiplets vertically or horizontally, optimizing each chiplet for its specific function, which enhances performance and energy efficiency.

**3. AI Capabilities**

**Q: What role does the Intel Neural Engine play in AI applications?**

**A:** It accelerates AI tasks directly on the processor, improving applications like machine learning, deep neural networks, and natural language processing.

**Q: What are the advantages of Neural Compute Engine (NCE) tiles?**

**A:** NCE tiles enable high computational throughput for deep learning applications, supporting both 16-bit floating point and 32-bit brain floating point formats.

**4. Integrated Graphics and Connectivity**

**Q: What advancements do Intel Arc Graphics provide?**

**A:** Intel Arc Graphics deliver high performance for gaming and multimedia tasks, offering significant improvements in image and video processing.

**8. Comparisons and Benchmarks**

**Q: Which connectivity technologies are supported by Intel Core Ultra processors?**  
**A:** These processors support Wi-Fi 6, Wi-Fi 7, and Thunderbolt 4, providing high-speed connectivity and reduced latency.

**5. Architecture and Internal Organization**

**Q: How does the Intel Thread Director optimize performance in hybrid architectures?**  
**A:** It dynamically assigns tasks to P-cores and E-cores based on workload requirements, ensuring energy efficiency and better task prioritization.

**Q: What is the significance of using Intel 4 manufacturing technology?**

**A:** Intel 4 enables higher transistor density and better performance, facilitating the integration of hybrid architecture and advanced features like AI and integrated graphics.

**6. Memory Hierarchy**

**Q: Describe the role of L1, L2, and L3 caches in the memory hierarchy.**  
**A:**

- L1 is the smallest and fastest cache for frequently accessed data.
- L2 is larger but slower and supports either single or multiple cores.
- L3 is the largest and slowest, shared across all cores to reduce latency in accessing the main memory.

**Q: How does the Intel Smart Cache benefit multi-core processing?**

**A:** It allows shared access among cores, enhancing data exchange and reducing memory access latency, particularly in multi-threaded workloads.

**7. Performance and Efficiency**

**Q: What metrics are used to evaluate processor performance?**  
**A:** Metrics include execution time, clock cycles per instruction (CPI), and benchmarks like Cinebench and 7-Zip performance.

**Q: How does Intel Adaptive Boost Technology improve processor performance?**

**A:** It dynamically increases core frequencies under favorable thermal and power conditions, optimizing performance for demanding tasks like gaming.

Diapo 17

ILP. This technique enables multiple instructions to execute simultaneously within a single CPU cycle.

Speculative Execution, allows the processor to predict and execute instructions before they are confirmed to be necessary.

Key features supporting both:

-Intel Thread Director which dynamically allocates workloads to the most suitable cores, either P or E cores.  
And the classics: BTB, GL predictors...

Present in previous intel architectures (no 100% seguro, se supone en base a que estaban antes)

Diapo 19

Cache Strategies, as mentioned before Intel Core Ultra processors use a hierarchical cache system  
-L1 dedicated to each core,  
-L2 cache shared diff with p-cores and e-cores  
-L3 shared across all cores

Dynamic Resource Allocation, The hybrid architecture, supported by the Intel Thread Director, along the integration of Intel Adaptive Boost for adjusting clock speeds.

Diapo 20

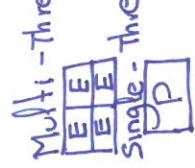
Lower Clock Cycles per Instruction (CPI) are achieved through parallel execution and optimized caching. The graph illustrates how the new architecture design enables higher clock frequencies with lower voltage.

Diapo 22

The Intel Neural Engine is a major innovation in Core Ultra processors compared to previous architectures. Unlike traditional reliance on external GPUs or TPUs, the Neural Engine brings on-chip AI processing. By integrating AI capabilities directly into the CPU, it accelerates workloads, reduces latency, and eliminates the need for separate accelerators in general-purpose applications.

Diapo 23

Intel 4 nodes are built thanks to Intel's advanced manufacturing process. These nodes are foundational to the Core Ultra processors, offering higher transistor density and greater energy efficiency. And here you can see the blueprint of the previous generation and the next ones that intel has in mind for the chips.



Diapo 24

And finally the Dynamic Workload Distribution as already mentioned before, using the Hybrid Architecture.  
P-cores focus on single-threaded tasks, delivering high performance for compute-intensive applications, while E-cores handle multi-threaded and background processes efficiently.

Diapo 25

Last but not least we will take a look at some benchmarks like this one, Cinebench 2024 that uses Redshift to benchmark CPU and GPU performance.  
-In multi-core performance, the Intel Core Ultra 9 285K leads with a score of 2523 using DDR5-8200 memory, surpassing competitors like the AMD Ryzen 9 9950X at 2345.  
-In single-core performance, the Intel Core Ultra 9 285K again takes the lead with a score of 150, followed closely by the Core Ultra 7 265K at 145.

Diapo 26

Cinebench 2024 power efficiency. The Intel Core Ultra 7 265K demonstrates impressive efficiency, delivering performance comparable to the AMD Ryzen 7 950X while consuming 29% less power than the Intel 14700K.  
However AMD is even more power-efficient, specifically the AMD's Ryzen 7 950X3D consuming 33% less while offering similar performance to the 265K.

Diapo 27

7-Zip benchmark evaluate CPU performance and memory efficiency using the LZMA (Lempel-Ziv-Markov chain algorithm) algorithm for compression and decompression tasks.

In file compression, the Intel Core Ultra 9 285K performs strongly, achieving 212 MIPS with DDR5-8200 memory, closely trailing AMD's top models like the Ryzen 9 9950X and 7950X. However, the Core Ultra 7 265K lags slightly, performing 6% slower than the Intel 14700K.  
In decompression tasks, the Core Ultra 9 285K remains competitive, scoring 203 MIPS, just below AMD's leading Ryzen 9 7950X3D at 205 MIPS. The 7 265K, however, shows its limitations without SMT support, falling 15% behind the 14700K.

- Compatibility with advanced technologies: Allows to integrate the Intel Thread Director

## INTEL ADAPTIVE BOOST

It is designed to maximize the performance of gaming and demanding applications while administrating energy limitations.

How:

1. Dynamical increase of frequencies of the cores
2. Advanced thermal management. Administers the digital sensors to avoid overheating

**SUPPORT DDR5 AND PCIe 5.0**  
**DDR5: → Better Memory Accesses, Out of Order scheduling**  
It improves speed and bandwidth allowing transfer up to 5600 MT/s  
Brings a flexible configuration which gives us compatibility with dual channels that improves the latency and the bandwidth. Including advanced technologies such as Data Scrambling and Out-Of-Order Scheduling.

## PCIe 5.0:

Gives us the ability of scalability and compatibility, this processor admits up to 20 PCIe entities. Allowing also the integration and flexibility with advanced configurations such as Lane Bifurcation.

## INTEGRATED GRAPHICS PROCESSING UNITS

Based on the Xe architecture and designed to offer a high graphic performance without a specific GPU.

Main characteristics:

- Scalar Xe Architecture
- Advance support for graphic APIs
- Accelerated codification and decodification by hardware
- Improvements in multimedia processing
- Higher video and games performance
- Low energetic cost

Diapo 11

## Internal Organization: Cores and Threads

Intel Core Ultra processors feature a hybrid architecture with multiple P-cores (high-performance) and E-cores (energy-efficient).

1. Caches: Three levels (L1, L2, L3) balance speed and size, reducing latency and enhancing performance.
2. Integrated Controllers: Built-in memory and I/O controllers streamline communication with memory and peripherals.
3. Security and Encryption: Technologies like Intel SGX protect data and applications. (Hablar de encriptación en Hardware sin tocar software)

4. AI Modules: Integrated AI modules that optimize tasks like natural language processing and image improve ..

Diapo 14

**Memory Hierarchy and Technology** optimizes data access by organizing memory from faster, smaller caches near the CPU to slower, larger ones further away. Intel Core Ultra processors include integrated memory controllers that handle data transfers and DRAM maintenance independently. Each controller manages half the physical memory address space in symmetric configurations, scaling bandwidth with increased channels to enhance efficiency.

## Cache Memory Levels:

- L1 Cache: Smallest, fastest, dedicated to individual cores for frequently accessed data, reducing latencies.
- L2 Cache: Larger, slower, shared or dedicated, storing reused task data efficiently.
- L3 Cache: Largest, slowest, shared across cores, bridging lower cache levels and RAM to reduce direct memory access.

Diapo 15

**Intel Smart Cache** serves as a shared, non-inclusive L3 cache for all AI cores and Processor Graphics. Adaptive Memory Technology tailors cache design for task-specific optimization:

- **P Cores:** Isolated L1/L2 caches per core for high-performance tasks, offering low latency.
- **E Cores:** Shared L2 caches across four cores, balancing energy efficiency and parallelism.

In both designs, L3 cache is shared across all physical cores, ensuring seamless data flow.

Diapo 16

Enseñar un poquito la idea de los Cores + Caches

### Diapo 3

- Intel biggest manufacturer of integrated circuits.
- Created the x86 architecture, all based in 1978 designs.
- Introduce in biggest part la Meteor Lake y el Foveros 3D ways of connecting and stacking components.
- Focus en la IA, GPU functions, more efficiency and performance.
- Also in security, with KeyLocker, that is encryption but in hardware, not touching software.

### Diapo 7

- Enhanced connectivity: Supports Wi-Fi 6, 7 and thunderbolt. High speed, reduced latency, more security and efficiency. Stable performance
- Target Market: High-end laptops and ultrabooks due to thermal solutions, exceptional performance and energy efficiency. For gaming, content creation and computational tasks

(De la 8 a la 10 leer sin más)

### Diapo 4

#### Diapo 10

### Diapo 4

#### CHARACTERISTICS OF THE ARCHITECTURE AND INTERNAL ORGANIZATION

One of the most important characteristics is the hybrid architecture  
If we talk about this hybrid architecture we have to differentiate between:

- P-cores: Are designed to carry out tasks that require high performance and they are optimized for intense workloads in only one thread. These cores are based on Lion Cove microarchitecture.
- E-cores: They produce a low energetic cost. They are designed for scalable multi-threaded performance, offloading background processes and low-priority tasks. They are based on Skymont microarchitecture.

#### INTEL THREAD DIRECTOR TECHNOLOGY

Improves the assignment of tasks in a system where we have to distribute the workload between the P-cores and E-cores. For this tasks we need to take into account some topics:

- Task prioritization: studies nature of the tasks
- Intelligent distribution: Based on the study
- Dynamical organization: Reorganizes loads depending on changing conditions
- Collaboration with the operating system

#### MANUFACTURING PROCESS

These processors of the PS series are manufactured using Intel 4, based in lithography of 10nm or smaller.

This brings different advantages for the process of fabrication:

- High density of transistors: Allowing the optimization of the space and increase performance.
- Energetic efficiency: Smaller transistors reduce the amount of energy required
- Support for new architectures: Support hybrid architecture.

Integrated graphics: Featured used specially for gaming and video processing.

- Efficiency and power management: Optimizes CPU, GPU and thermal performance for portable devices with Adaptive Boost. Better hardware life..
- Memory and IOs: DDR5 and LPDDR5 supported, the second one for mobiles and laptops

**Q: How does Intel Core Ultra compare to previous architectures in benchmarks?**

**A:** Intel Core Ultra processors outperform previous architectures in both single-core and

multi-core tasks, with improvements in energy efficiency and AI integration.

**Q: What makes Intel Core Ultra processors suitable for gaming and content creation?**

**A:** Their hybrid architecture, advanced AI capabilities, integrated Intel Arc Graphics, and dynamic resource management make them ideal for high-performance applications.

Feel free to use or adapt these questions during an explanation session!



Final written exam (January 2022)      Test      Year 2021/22

Name: \_\_\_\_\_

Aula: \_\_\_\_\_ Mesa número: \_\_\_\_\_

Indicate the correct answer to each test question in the next table:

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23	24	25	26	27	28	29	30

**Important:** Test counts **3 points**. Each correct question counts 0.1 points. Wrong answers subtract 0.1 points. Not answered questions are not considered for the mark.

1. A current smartphone can exploit:  
 a) Thread-level parallelism  
 b) Instruction-level parallelism  
 c) CPU-level parallelism  
 d) All the above
2. According to the Flynn's taxonomy, GPUs are:  
 a) SISD computers  
 b) MISD computers  
 c) SIMD computers  
 d) MIMD computers
3. We want to analyze the performance of two computers, M1 and M2, which support the same instruction set. In this instruction set there are three types of instructions, namely, A, B, and C. M1 needs 3, 6, and 4 clock cycles to execute instructions of type A, B, and C, respectively. M2 needs 4 cycles for all types. M1 runs at 400 MHz and M2 runs at 750 MHz. The test program we will use has the following mix of instructions: 40% of type A, 50% of type B, and 10% of type C. Indicate the peak speed in MIPS for M1 and M2.  
 a) 133.3 for M1 and 187.5 for M2  
 b) 87 for M1 and 187.5 for M2  
 c) 125 for M1 and 155 for M2  
 d) None of the above
4. If the hard disk of a computer is replaced by another disk six times faster than the current one, calculate the use that should be made of this device if we want to make the computer run three times faster  
 a) It is not possible  
 b) The disk must be used 20% of the time  
 c) The disk must be used 80% of the time  
 d) None of the above
5. Consider an iteration of the next loop:  
 LOOP: LW R1, 0(R2) ; I1  
 DADD R1, R1, R3 ; I2  
 SW R1, 0(R2) ; I3  
 DADDUI R2, R2, #4 ; I4  
 DADDUI R4, R4, #-1 ; I5  
 BNEZ R4, LOOP ; I6
- a) There is a data dependence between I3 and I4  
 b) There is an antidependence between I3 and I4
6. Consider two consecutive iterations of the following loop:  
 LOOP: LW R1, 0(R2) ; I1  
 DADD R1, R1, R3 ; I2  
 SW R1, 0(R2) ; I3  
 DADDUI R2, R2, #4 ; I4  
 DADDUI R4, R4, #-1 ; I5  
 BNEZ R4, LOOP ; I6
- a) There is a data dependence between I5 and I6  
 b) There is an antidependence between I1 and I2  
 c) There is an antidependence between I1 and I3  
 d) There is an output dependence between I1 and I3
7. In a pipelined processor implementing forwarding and using pipelined FP functional units, a latency of 3 cycles between a FP addition and an instruction using the result of this operation implies that...  
 a) The instruction using the result must be stopped for 2 cycles  
 b) The instruction using the result must be stopped for 3 cycles  
 c) The execution step of the FP addition takes 2 cycles  
 d) The execution step of the FP addition takes 3 cycles
8. An antidependence can lead to a:  
 a) WAW hazard  
 b) RAW hazard  
 c) WAR hazard  
 d) None of the above
9. Indicate the correct option:  
 a) The Scoreboard technique solves RAW and WAR hazards  
 b) The Scoreboard technique solves WAR and WAR hazards  
 c) The Tomasulo technique (without speculation) solves RAW and WAR hazards  
 d) The Tomasulo technique (without speculation) solves WAR and WAR hazards
10. Software pipelining.  
 a) Is implemented by the compiler to reduce stalls due to WAR hazards  
 b) Is implemented by the CPU to reduce stalls due to WAR hazards  
 c) Is implemented by the compiler and the CPU to reduce stalls due to WAR hazards  
 d) None of the above
11. The loop unrolling technique:  
 a) Is a dynamic technique to schedule the instructions and increase the ILP  
 b) Is a static technique to schedule the instructions and reduce the size of the object code  
 c) Is a static speculative technique to schedule the instructions and increase the ILP  
 d) Is a static technique to schedule the instructions and reduce the size of the object code
12. We have a one level 2-bit branch predictor initialized to strongly not taken. Indicate the correct sequence of hits (H) and misses (M) for B1 in the next code:  
 start: DADDIU R2, R0, #6  
 DADDIU R4, R0, #1  
 loop: DADDIU R2, R2, #-1  
 BNEZ R2, loop ; B1  
 DADDIU R4, R4, #-1  
 a) HHMMMM  
 b) HHHHHH  
 c) MMHHHH  
 d) MMHHHH

13. Indicate the width and the number of bits to implement an implicit BTB with 1024 entries, knowing that the address bus uses 24 bits (there are several solutions, but only one of the options is correct).

- a) Width = 24, Number of bits =  $1024 \times 24$  bits
- b) Width = 1024, Number of bits =  $1024 \times 1024$  bits
- c) Width =  $24 \times 24$ , Number of bits =  $1024 \times 24 \times 24$  bits
- d) Width = 24, Number of bits =  $2^{24} \times 24$

14. What is the value of a 4-input BHT knowing that we have just executed a loop 100 times and now continue executing the following instructions sequentially?

- a) 1111
- b) 1110
- c) 0101
- d) 0011

15. Indicate the correct option for the multipath execution technique:

- a) Both paths of a conditional branch are executed, and the instruction is executed depending on the value of a predicate register
- b) One of the paths of a conditional branch is chosen to be executed by using prediction
- c) Both paths of a conditional branch are executed, without using prediction
- d) One of the paths of a conditional branch is randomly chosen to be executed

16. Software speculation...

- a) Is supported by hardware (dynamic) branch prediction
- b) Speculative instructions resulting in terminating exceptions always produce program termination
- c) Is based on the Scoreboard algorithm
- d) None of the above

17. In the poison bits technique, if a normal (non-speculative) instruction results in a terminating exception:

- a) The execution of the program continues, but the poison bit of its destination register is set
- b) The execution of the program continues, but the poison bit of its destination register is reset
- c) The execution of the program is stopped
- d) None of the above

18. In the *commit* step of the speculative version of Tomasulo ...

- a) Instruction results are provided to other instructions that need them
- b) The ROB is employed to complete the execution of instructions according to the program order
- c) If the instruction to be completed is a branch, the destination register (or the memory) is updated
- d) None of the above

19. A grade 3-way superscalar has an ideal CPI of

- a) 3
- b) 1.3
- c) 1/3
- d) 1

20. Regarding the branch prediction in a superscalar:

- a) It is carried out in the fetch stage.
- b) It is carried out in the decode stage.
- c) It is carried out in the issue stage.
- d) It is carried out in the commit stage.

21. One of the drawbacks of superscalars is:

- a) They are not compatible at the binary level with a scalar machine.
- b) They introduce many NOP operations in the pipeline.
- c) The hardware is so simple that they have problems due to branches.
- d) There is a higher probability of risks in instructions that are launched at the same time.

22. MPI is:

- a) A software product sold by Microsoft
- b) A message-passing library specification
- c) A compiler directive standard
- d) All the above

- a) Since they use static scheduling, they are highly sensitive to cache misses.
- b) They have a register bank that is complex in area and access time.
- c) They duplicate hardware resources to be more efficient.
- d) If the compiler is not efficient the size of the machine code can be exceptionally large.

- a) Distributed memory computers
- b) Shared memory computers with uniform access time
- c) Shared memory computers with non-uniform memory access
- d) Distributed cache memory computers

- a) An evolution of the P6 microarchitecture aimed at increasing the number of instructions to be processed per cycle
- b) Increasing clock frequency
- c) The P6 microarchitecture, keeping the decoder structure
- d) The Pentium Pro architecture, but with Xeon improvements

- a) In the heterogeneous computing paradigm:
  - a) The CPU performs general purpose computing
  - b) Computing systems become a hybrid structure
  - c) Specific applications are run on accelerators
  - d) All the above
- b) On multicore processors with HyperThreading, better execution times for two applications are achieved by using:
  - a) The same core for both applications
  - b) A different core for each application
  - c) Depends on the number of threads per core allowed by the system
  - d) None of the above
- c) 27. A NUMA system:
  - a) Uses a shared memory
  - b) Uses the Internet to connect its processors
  - c) Has a variable memory access latency
  - d) None of the above
- d) 28. Indicate the correct statement:
  - a) A SIMD architecture is less energy-efficient than a MIMD architecture
  - b) The MMX extensions on Intel processors are MISD extensions
  - c) There is no penalty due to context switching in GPUs
  - d) None of the above
- e) 29. Skype or Gmail are examples of:
  - a) SaaS Cloud
  - b) PaaS Cloud
  - c) IaaS Cloud
  - d) JaaS Cloud
- f) 30. MPI is:
  - a) A software product sold by Microsoft
  - b) A message-passing library specification
  - c) A compiler directive standard
  - d) All the above

LD	F1, O(R2)	IF ID EX ME WB
LD	F2, O(R2)	IF ID EX ME WB
MULD	F1, F1, F10	EX IF ID EX ME WB
MULD	F2, F2, F20	EX IF ID EX ME WB
ADDO	F3, F1, F2	EX IF ID EX ME WB
DIVD	F3, F3, F30	EX IF ID EX ME WB
SD	F3, O, (R3)	EX IF ID EX ME WB



Name: Important: Solve this exercise on this sheet. You can use more sheets if you need, but do not forget to write down your name in all the sheets you deliver.

**Exercise 3 (2.5 points).** Consider a 4-way superscalar processor capable of issuing four instructions of any type per cycle, with five functional units (F1, F2, F3, F4, and F5). Assume that instructions require one cycle for the decode stage, one cycle for the execute stage and one cycle for the write stage. The result of an instruction is available to other instructions at the next cycle in which that result is written.

Given the next code:

```
LW R1, 0(R2)      # I1
DADD R4, R1, R3   # I2
DSUB R6, R7, R8   # I3
LW R3, 0(R9)      # I4
DSUB R10, R10, R12 # I5
DADD R8, R8, R1    # I6
DSUB R3, R3, R1    # I7
```

Draw a diagram showing, for each clock cycle, which instructions are in the decode step (four at most), in the execute step (five at most), and in the write step (four at most), for each of the following situations:

- In-order execution and in-order writing.
- Out-of-order execution, without renaming capacity, and in-order writing.
- Out-of-order execution with renaming (assuming enough registers for this), and out-of-order writing.

Use the following tables to provide your solutions, detailing the instruction status at each cycle.

- In-order execution and in-order writing.

Cycle	Decode	Execute	Write
1	I1		
2			
3			
4			
5			
6			
7			
8			
9			

- Out-of-order execution, without renaming capacity, and in-order writing.

Cycle	Decode	Execute	Write
1	I1		
2	I2		
3	I3		
4			
5			
6			
7			
8			
9			

- c) Out-of-order execution with renaming (assuming enough registers for this), and out-of-order writing.

Cycle	Decode	Execute	Write
1	I1	I2	I4
2	I5	I6	I7
3	I7	I8	I9
4			
5			
6			
7			
8			
9			

**A** Dos instrucciones no pueden leer el mismo registro a la vez  
**A** Cada ciclo se escribe en un solo registro antes de la vec de que se uses

# No hay que esperar que se libere la FU

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING COMPUTING SYSTEMS DEPARTMENT COMPUTER SCIENCE DEGREE. COMPUTER ARCHITECTURE	Exercises	Year 2021/22
Final written exam (January 2022)		

Name: \_\_\_\_\_

**Important:** Solve this exercise **on this sheet**. You can use **more sheets** if you need, but do not forget to write down your name in all the sheets you deliver.

**Exercise 1 (2.5 points).** The next loop is executed on a MIPS64 pipelined processor to compute the weighted mean of the marks obtained by a group of Computer Architecture students in two exam exercises. Assume that these marks have been previously stored (as double-precision floating-point values) from the memory addresses contained in R1 (exercise 1 marks) and R2 (exercise 2 marks) and that the resulting mean values must be stored from the memory address contained in R3, R4 contains the number of student marks, F10 and F20 contain the weight assigned to each exercise, and F30 has been initialized to the sum of the weights.

```

Loop: L.D F1, 0(R1) // gets the mark from exercise 1
      L.D F2, 0(R2) // gets the mark from exercise 2
      MUL.D F1, F1, F10 // weights the mark from exercise 1
      MUL.D F2, F2, F20 // weights the mark from exercise 2
      ADD.D F3, F1, F2 // obtains the weighted mean value
      DIV.D F3, F3, F30 // stores the result
      S.D F3, 0(R3) // points to the next mark from exercise 1
      DADD.D R1, R1, #8 // points to the next mark from exercise 2
      DADD.D R2, R2, #8 // points to the next mark from exercise 2
      DADD.D R3, R3, #8 // points to the next weighted mean
      DADD.D R4, R4, #-1
      BNEZ R4, Loop // more student marks?
  
```

The processor exhibits the usual characteristics: a value can be written and read from a register in the same clock cycle, dependences are detected and solved in the ID stage, operand forwarding is supported, branches are completely solved in the ID stage, the integer functional unit requires 1 cycle, and the execution step of L.D and S.D requires 1 cycle. Functional units for floating-point addition, multiplication, and division are fully pipelined, and require, respectively, 2, 3, and 5 cycles.

- Detail the stalls (be careful! NOT DEPENDENCES) required to execute each iteration of the original loop and compute the CPR value.
- Unroll the loop the minimum number of times to schedule it without stalls. Assume that delayed branch is supported. Keep in mind that the program behavior must be preserved.
- Compute the CPR value for the resulting loop (after applying loop unrolling and scheduling).

**NOTE:** The loop computes the weighted mean for each pair of student marks ( $x_1$  and  $x_2$ ) by using the expression  $\bar{x} = \frac{x_1 \cdot w_1 + x_2 \cdot w_2}{w_1 + w_2}$ , where  $w_1$  and  $w_2$  are the weights considered for exercise 1 and exercise 2.

*LD EX ME W1  
F INDEX ME W2  
IF ID IF ID IF ID*

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING COMPUTING SYSTEMS DEPARTMENT COMPUTER SCIENCE DEGREE. COMPUTER ARCHITECTURE	Exercises	Year 2021/22
Final written exam (January 2022)		

Name: \_\_\_\_\_

**Important:** Solve this exercise **on this sheet**. You can use **more sheets** if you need, but do not forget to write down your name in all the sheets you deliver.

**Exercise 2 (2 points).** We assume a MIPS64 processor using hardware speculation and dynamic scheduling using the Tomasulo technique with speculation, with the following characteristics:

- The load and store execution stage takes 1 cycle.
- The floating-point multiplication (FP) execution step takes 6 cycles.
- The floating-point addition (FP) execution step takes 4 cycles.
- A 4-input reorder buffer (ROB) is available.
- It has a 2-entry reservation stations for each unit in FP.
- It has a 2-entry load/stage input load buffer.

- The execution stage of an instruction that is waiting for a result cannot start in the same cycle in which the result is written to the CDB (Common Data Bus).
  - A resource that is released in one clock cycle cannot be used until the next clock cycle.
- a) Detail the execution of the code below when the basic version of Tomasulo's algorithm with speculation is used. To do so, indicate in the table below the start and end cycle number for each of the steps of each instruction (NB: For 1-cycle instruction steps, you must annotate just 1 cycle in the table). (5)

Instruction	Issue	Execution	Write result	Commit
MUL.D F4,F6,F12	1	2-7	8	+1 (*)
SUB.D F6,F4,F10	2	9-12	13	+1 (*)
L.D F4,10(R2)	3	4	5	+2 No ROB space
DIV.D F0,F14,F16	4	8-13	14	+4
SUB.D F12,F14,F0	5	14-17	15	-
ADD.D F20,F28,F24	6	15-18	19	-
L.D F0,10(R2)	7	19-22	23	-
MUL.D F18,F28,F24	8	23	24	-
	9	24	26	-

- b) Complete the table below to display the contents of the ROB (which has only 4 entries), just in the same cycle in which the first L.D instruction ends. (5)

Entry	Busy (yes/no)	Instruction	Status	Destination	Value
0	YES	HULD4,F6,F12	EXEC	F4	F6+F12
1	YES	SUBD F6,F4,F10	Wait Issue	F6	F4-F10
2	YES	LD F9,10(R2)	WRITE	F9	MEM[10]
3	YES	DIVDE0,F14,F16	WRITE	F16	F14*F0

*LD LD Shal MUL MUL*

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING	 UCLM
COMPUTING SYSTEMS DEPARTMENT	
COMPUTER SCIENCE DEGREE. COMPUTER ARCHITECTURE	

Final written exam (January 2023) Year 2022/23

Name: \_\_\_\_\_

Indicate the correct answer to each test question in the next table:

A	B	C	D	B	C	D	B	A	C	D	A	B	C	D
16	17	18	19	20	21	22	23	24	25	26	27	28	29	30

**Important:** Test counts **3 points**. Each correct question counts 0.1 points. Wrong answers subtract 0.13 points. Not answered questions are not considered for the mark.

1. Synthetic benchmarks programs:

- a. They are specially designed to measure the performance of an individual component of a computer.
- b. They are specially designed to measure the performance of an entire computer.
- c. They are specially designed to measure the speed of access to a database system.
- d. None of the above.

2. The organization of a computer should be designed to:

- a. Survive for many years.
- b. Implement a particular architecture specification. → Ex : x86
- c. Implement several architectures.
- d. None of the above.

3. If a machine A executes a program in 30 seconds and a machine B executes the same program in 45 seconds, how much faster is A than B?

- a. A is 50% faster than B
- b. A is 150% faster than B
- c. B is 1.5 times faster than A
- d. B is 50% faster than A

4. Suppose a program takes 200 seconds to run on a machine, 150 seconds of which are for multiplication operations. How much do I have to improve the multiplication speed if I want the program to run twice as fast?

- a. 1.5 times
- b. 2 times
- c. 3 times
- d. None of the above

5. In a pipelined MIPS processor supporting forwarding, when executing two consecutive integer instructions:

- a. It is never necessary to stop the pipeline.
- b. It is necessary to stop if both instructions are arithmetic and present a data dependence.
- c. If the first instruction is a memory access, sometimes it is necessary to stop.
- d. None of the above.

6. In the following loop:  
**Loop :** L.D F8, 0(R1) ;I1  
ADD.D F6, F2, F8 ;I2  
S.D F6, 0(R1) ;I3  
DADDUI R1, R1, #-8 ;I4  
BNEZ R1, Loop ;I5

Between instructions I2 and I3...

- a. There is an antidependence between iterations
- b. There is an output dependence in the same iteration
- c. There is a data dependence between iterations
- d. There is an antidependence in the same iteration

7. The loop unrolling technique pursues... → even to branches
- a. Reduce stalls due to W/A hazards
  - b. Reduce stalls due to W/W hazards
  - c. Reduce stalls due to structural hazards
  - d. Reduce stalls due to control hazards
8. In Scoreboard, instructions always get their operands:
- a. Through the common data bus (CDB)
  - b. Through the register file (or register bank)
  - c. By means of the forwarding technique
  - d. From the reorder buffer (ROB)
9. In the Scoreboard technique:
- a. WAW and WAR hazards cannot appear
  - b. WAW and WAR hazards may appear, but they are not controlled
  - c. WAW and WAR hazards are checked in the first and last step, respectively
  - d. None of the above. → Scoreboard check
10. The Qi and Qk files are in a reservation station in Tomasulo store:
- a. The value of the source operands needed by the corresponding instruction
  - b. The reservation stations that will produce the instruction source operands.
  - c. The registers that contain the instruction source operands.
  - d. None of the above.
11. What is the latency of a branch instruction?
- a. The number of additional cycles between the IF stage and the cycle in which the branch is resolved.
  - b. The number of additional cycles between the ID stage and the cycle in which the branch is resolved.
  - c. The number of cycles between the IF and ID stages.
  - d. The number of cycles between the IF and EX stages.
12. Assuming a conditional branch, the miss penalty in the prediction depends on
- a. The stage where the branch target address is known.
  - b. The stage where the instruction source operands.
  - c. a) and b) are true.
  - d. All the above are false.
13. For a 5-stage pipeline, we run a code with 30% of instructions being a conditional branch, from which 80% are taken. In this processor, the branch direction is solved in the EX stage and the target address in the ID stage (the CPI is 1 for the rest of the instructions). Indicate the correct answer:
- a. CPI = 1 + 3 x 0.8 x 0.3
  - b. CPI = 1 + 2 x 0.8 x 0.3
  - c. CPI = 1 + 3 x 0.2 x 0.7
  - d. CPI = 1 + 2 x 0.2 x 0.7
14. Apart from storing prediction bits, what is the purpose of the branch target buffer (BTB) in the computer architecture?
- a. To store the branch target addresses of branch instructions.
  - b. To reduce the latency cycles of branch instructions.
  - c. Exception management.
  - d. All are true.

15. For the software speculation (without hardware support), which of the following sentences is false?
- The compiler bets for a specific branch direction, moving instructions of that block to the former block.
  - It is required to perform an efficient branch (static) prediction during the compiling time.
  - During the execution of a program, the processor bets for a branch direction and the instructions of this block are executed speculatively. These instructions will be valid when the branch result is finally known.
  - Renaming is used to allow the speculative execution of some instructions without modifying the data flow.
16. Let's assume we use software speculation when running a code. In the case that a speculative instruction causes a terminating exception, and the software speculation is not correct, we can say that:
- If we use the poison bits technique, the exception will always take place and the program will terminate in error.
  - If we use the poison bits technique, the exception will only take place in some cases in which the program will terminate in error.
  - Regardless the technique used to deal with exceptions, the exception will always take place and the program will terminate in error.
  - None of the above.
17. In a superscalar, in the issue stage ...
- Name dependencies are removed by renaming.
  - Instructions are read from the reorder buffer and sent to execution.
  - An arbiter selects which instructions can be sent to execution.
  - All are true.
18. A characteristic of dynamic superscalar processors is that they always ...
- Use dynamic scheduling.
  - The number of instructions issued per cycle is fixed.
  - Instructions are executed in order.
  - All are true.
19. We have a VLIW processor that can issue four instructions per cycle. A compiler receives as input a program with 36 instructions and provides as output a program for this VLIW processor with an average efficiency of 3 operations per cycle. We can state that...
- 3 long instructions at the compiler output contain 4 NOP operations.
  - 4 long instructions at the compiler output contain 3 NOP operations.
  - The minimum number of long instructions in the compiler output is 9.
  - The minimum number of long instructions in the compiler output is 12.
20. Through the execution of a test program, we have obtained that the used superscalar has an ideal CPI of 0.25. What can you deduce from this statement?
- The processor uses out-of-order execution but does not use speculation.
  - The processor uses in-order execution and speculation.
  - The processor has a frequency of  $\frac{1}{4}$  MHz
  - The processor is a 4-way superscalar.
21. In a superscalar processor pre-decoding is performed:
- At the issue stage
  - Between L2 and L1
  - At the write-back stage
  - None of the above is correct
22. Each iteration of a given loop in a scalar processor requires 10 clock cycles to be issued. We unrolled such a loop 8 times and scheduled it for a 4-way superscalar obtaining that we now perform 8 iterations in 10 cycles. What is the speedup obtained?
- 4
  - 6
  - 8
  - 10
23. In the period 2004 to 2020:
- The number of transistors on a chip has increased exponentially.
  - The working frequency of processors has remained almost stable.
  - It has slightly increased the performance of a single thread.
  - All are true
24. The architecture that is open source is:
- The x86 family
  - ARM
  - RISC-V
  - All are false
25. The market sector where Intel processors are more dominant with respect to AMD is:
- Laptops
  - Desktop
  - Servers
  - IoT
26. Cortex-M are ARM processors whose application scope is:
- Low power consumption for microcontrollers and IoT.
  - Real-time for critical applications
  - HPC
  - Applications for computationally complex systems
27. A machine with tensor processors is an architecture according to Flynn's classification:
- SISD
  - SIMD
  - MISD
  - MIMD
28. CUDA:
- It is C language with minimal extensions
  - Uses a multitude of processing units and a single control unit
  - Defines a programming model
  - All are true
29. Which of the following statements is true?
- CLOUD is a collection of computers owned by multiple owners in different locations and interconnected so that users can share the combined power of resources.
  - GRID is a collection of computers usually owned by a single party.
  - Facebook is an example of CLOUD defined as SaaS (Software as a Service).
  - None is true
30. A standard that by means of directives tells the compiler the region of the program to parallelize is
- OpenMP
  - PVM
  - MPI
  - None is true

Name: \_\_\_\_\_

**Important:** Solve each exercise in a different sheet. You can use more sheets if you need, but do not forget to write down your name in all the sheets you deliver.

**Exercise 1 (2.5 points).** Let's consider the following code in MIPS64. You need to assume that registers F2, F4, F6 and F8 are firstly initialized to 0, and that registers RA, RB and RC contain valid memory address.

```

    ADDI R1, R0, #0
    ADDI R2, R0, #3
    SLT R3, R1, R2
    BEQ R3, R0, END_L1
    L.D F2, 0(RA)
    L.D F4, 0(RB)
    ADDI R4, R0, #4
    ADDI R5, R0, #4
    SLT R6, R4, R5
    BEQ R6, R0, END_L2
    MUL_D F6, F2, F4
    ADD_D F8, F8, F6
    ADDI R4, R4, #1
    J LOOP2
    S.D F8, 0(RC)
    ADDI R1, R1, #1
    ADDI RA, RA, #8
    ADDI RB, RB, #8
    ADDI RC, RC, #8
    J LOOP1
    /* Fin del programa */

You need to perform the following activities:
```

- a) Fill the following table, bearing in mind the actual behavior of conditional Branch instructions B1 and B2:

Step	Executed branch instruction (B1 or B2)	R1 value	R4 value	Actual behavior of branch instruction (Taken/Not Taken)
1	B1	0	0	NT
2	B2	0	0	NT

(Use as many files as you need in a separate sheet)

- b) Calculate the hit predictions rate for a 2-bit predictor, which is initialized to weakly not taken (WNT = 0). 1) Use the same table of section "a" and add the corresponding columns to show the predictions at every step.  
 2) Calculate the hit predictions rate for a (2,2)-correlated predictor. All the data structures for this predictor are initialized to 0. Use a table like that of section "a" and show the values for data in these structures for every step in the table.

- d) Which predictor achieves the best performance for the assembly code in this exercise? Please, you need to justify your response.

Name: \_\_\_\_\_

**Important:** Solve each exercise in a different sheet. You can use more sheets if you need, but do not forget to write down your name in all the sheets you deliver.

**Exercise 2 (2 points).** We assume a MIPS64 processor using hardware speculation according to the Tomasulo technique, with the following characteristics:

- The execution step of loads and stores takes 2 cycles.
- There are three fully pipelined Floating-Point (FP) functional units:
  - A FP adder (also available for subtracting) that requires 4 cycles,
  - A FP multiplier that requires 6 cycles, and
  - A FP divider that requires 12 cycles.
- A 4-input reorder buffer (ROB) is available.
- There is a 2-entry reservation station per FP unit.
- There is a 1-entry load buffer and a 1-entry store buffer.
- The execution stage of an instruction that is waiting for a result cannot start in the same cycle in which the result is written to the CDB (Common Data Bus). → Se pilla do carryon ↗
  - A resource (the ROB, a reservation station, etc.) that is released in one clock cycle cannot be used until the next clock cycle.

- a) Detail the execution of the code below. To do so, indicate in the table below the start and end cycle number for each of the steps of each instruction (1 point), and, when applicable, the reason to stall the execution of the corresponding step (a RAW hazard, a structural hazard, the program order, etc.) (0.5 points).

Instruction	Issue cycles(s)	Execution cycles(s)	Write result cycle(s)	Commit cycle(s)	Stall cause
MUL_D F4, F6, F2	1	2-7	8	9	
SUB_D F6, F4, F10	2	9-12	13	14	
L.D F4, 0(R1)	3	4-5	6	15	ORDER
DIV_D F0, F4, F6	4	19-25	26	27	ORDER
L.D F8, 0(R2)	5	16	17	28	
SUB_D F2, F8, F0	6	17	27-30	31	
ADD_D F0, F8, F2	7	18	27-35	36	
S.D F0, 0(R1)	8	28	27-38	39	

- b) Detail the contents of the ROB during the "Issue" cycle of SUB\_D F2, F8, F0 (0.5 points).

Entry number	Busy (Y/N)	Instruction	State	Destination	Value
0	Y	LD F8, 0(R2)	exec	F8	MEM_R2→F8
1	Y	SUB_D F2, F8, F0	Issue	F2	F8-F10
2	-		-	-	-
3	Y	DIV_D F0, F4, F6	exec	F0	F4/F6

Cuanto más esté, más bajo se va el puntaje.



# Ejercicios Indio

1

Loop: LD F2,0(R1)

MULD F4,F2,F0

MULD F10,F4,F8

LD F6,0(R2)

ADD.D F10,F10,F8

SD F10,0(R2)

DADDIU R1,R1,8

DADDIU R2,R2,8

CMPI R1,800

BEQZ R1,loop

VLIW

- 2 MEM references

- 1 FP op

- 1 Int op / branch

FP → FP	Latency
FP → Store	3
LD → FP	2
LD → Store	1
	0

Fixed → F0, F8

a) Unroll 4 times to no stall

LD F2,0(R1)

LD F12,8(R1)

LD F22,16(R1)

LD F1,24(R1)

DADDIU R1,R1,32

MULD F4,F2,F0

MULD F14,F12,F0

MULD 24,F22,F0

MULD F3,F1,F0

DADDIU R2,R2,32

MULD F10,F4,F8

MULD F20,F14,F8

MULD F30,F24,F8

MULD F5,F3,F8

CMPI R1,800

LD F6,0(R2)

LD F16,0(R2)

LD F26,0(R2)

LD F7,0(R2)

BEQZ R1,loop

ADD.D F10,F10,F8

ADD.D F20,F20,F8

ADD.D F30,F30,F8

ADD.D F9,F9,F8

SD F10,0(R2)

SD F20,8(R2)

SD F30,16(R2)

SD F9,24(R2)

Mem 1	Mem 2	FP	Int
LD F2,0(R1)	LD F6,0(R2)		
LD F12,8(R1)	LD F16,8(R2)		
LD F22,16(R1)	LD F26,16(R2)	MULD F4,F2,F0	
LD F1,24(R1)	LD F7,24(R2)	MULD F14,F12,F0 MULD F24,F22,F0	
		MULD F3,F1,F0	
		MULD F10,F4,F8	
		MULD F20,F14,F8	
		MULD F30,F24,F8	
		MULD F5,F3,F8	
		ADDD F10,F10,F8	
		ADDD F20,F20,F8	
		ADDD F30,F30,F8	
SD F10,0(R2)		ADDD F9,F9,F8	DADDIU R1,R1,32
SD F20,8(R2)			DADDIU R2,R2,32
SD F30,16(R2)			CMPI R1,800
SD F9,24(R2)			BEQZ R1,loop
-8			

↳ ¿Pq? → sobrepasa el +32 de hace dos ciclos

→ Explicación. Hay dependencia. Cuando un LD crea un dato y un FP lo necesita, hay stall de 1. Por eso el hueco OBLIGATORIO entre ambas

loop: LD F0,0(R4)

ADDD F4,F0,F2

SD F4,0(R1)

DADDIU R1,R1,-8

BNE R1,R2,loop

Latency cycles

FP → FP	3
FP → store	2
LD → FP	1
LD → store	0

## Software Pipelining

loop: LD F0,0(R1)  
 ADD.D F4,F0,F2  
 SD F4,0(R1)  
 DADDUI R1,R1,-8  
 BNE R1,R2,loop

Iteration 1	Iteration 2	Iteration 3
LD F0,0(R1)	LD F0,-8(R1)	
ADD.D F4,F0,F2	ADD.D F4,F0,F2	LD F0,-16(R1)
SD F4,0(R1)	SD F4,-8(R1)	ADD.D F4,F0,F2
		SD F4,-16(R1)

→ Modified loop

bop: SD F4,0(R1)  
 ADD.D F4,F0,F2  
 LD F0,-16(R1)  
 DADDUI R1,R1,-8  
 BNE R1,R2,loop

{ DIFF iterations

Start-up:

LD F0,0(R1)  
 ADD.D F4,F0,F2  
 LD F0,-8(R1)  
 DADDUI R2,R2,16

loop: LD R4,0(R1)  
 LD RS,0(R2)  
 SLT R4,R4,RS  
 SD R4,0(R1)  
 DADDUI R1,R1,8  
 DADDUI R2,R2,8  
 BNE R1,R3,loop

Start-up:

LD R4,0(R1)  
 LD RS,0(R2)  
 SLT R4,R4,RS  
 LD R4,8(R1)  
 LD RS,8(R2)  
 DADDUI R3,R3,-16

{ independent iteration 1

LD R4,0(R1)  
 LD RS,0(R2)  
 SLT R4,R4,RS  
 LD R4,8(R1)  
 LD RS,8(R2)  
 SD R4,0(R1)      SLT R4,R4,RS  
 SD R4,8(R1)

LD.R4,16(R1)  
 LD RS,16(R2)  
 SLT R4,R4,RS  
 SD R4,16(R1)

loop:

SD R4,0(R1)  
 SLT R4,R4,RS  
 LD R4,16(R1)  
 LD RS,16(R2)  
 DADDUI R2,R1,8  
 DADDUI R2,R2,8  
 BNE R1,R3,loop

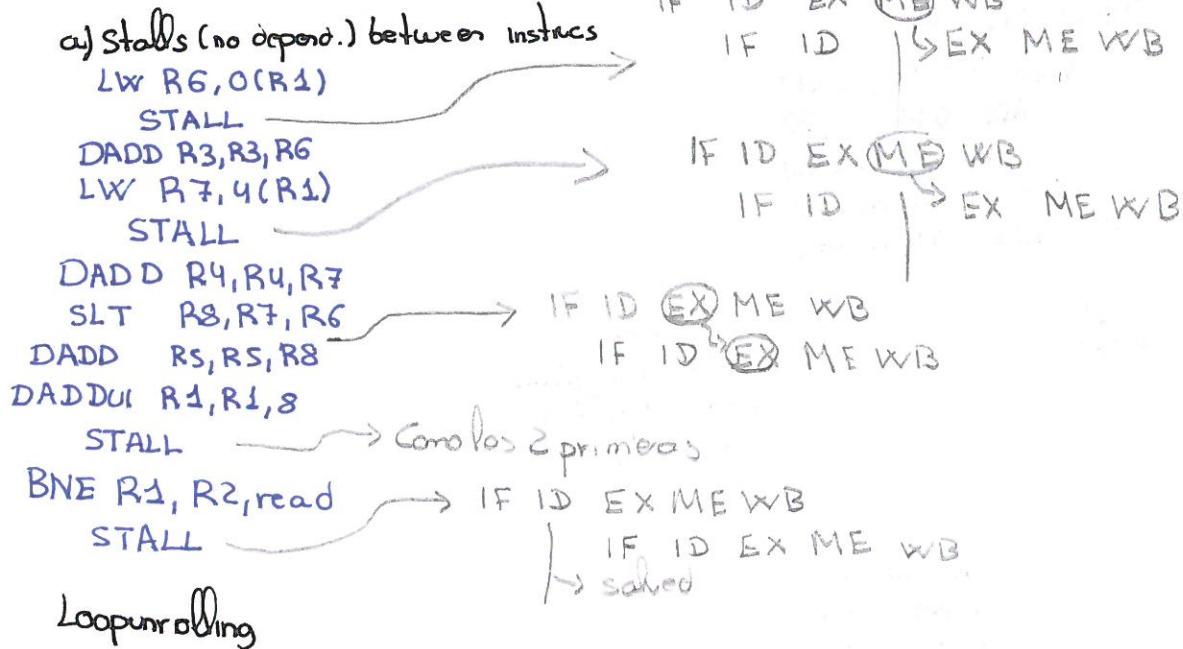
finish-up:  
 SDR4,0(R1)  
 SLT R4,R4,RS  
 SD R4,8(R1)

↓ al contrario,  
 y 0,-8,-16 por  
 iteración

read: LW R6,0(R1)  
 DADD R3,R3,R6  
 LW R7,4(R1)  
 DADD R4,R4,R7  
 SLT R8,R7,R6  
 ADD RS,RS,R8

DADDUI R1,R1,8  
 BNE R1,R2,read

- Forwarding ✓
- Write and read at same cycle
- Dependencies and branches solved at ID
- Delayed branch



read: LW R6,0(R1)  
 DADD R3,R3,R6  
 LW R7,4(R1)  
 DADD R4,R4,R7  
 SLT R8,R7,R6  
 DADD RS,RS,R8  
 LW R9,8(R1)  
 DADD R3,R3,R9  
 LW R10,12(R1)  
 DADD R4,R4,R10  
 SLT R11,R10,R9  
 DADDUI R1,R1,16  
 BNE R1,R2,read

LW R6,0(R1)  
 LW R7,4(R1)  
 LW R9,8(R1)  
 LW R10,12(R1)  
 ⇒ DADD R3,R3,R6  
 DADD R4,R4,R7  
 DADD R3,R3,R9  
 DADD R4,R4,R10  
 SLT R8,R7,R9  
 DADD RS,RS,R8  
 SLT R11,R10,R9  
 BNE R1,R2,read  
 DADD RS,RS,R11

{ Renaming and  
hence stalls

LD	F1, O(R1)	IF	ID	EX	ME WB
LD	F2, O(R2)	IF	ID	EX	ME WB
MULD	F1, F1, F10	IF	ID	EX	ME WB
MULD	F2, F2, F20	IF	ID	EX	ME WB
ADDD	F3, F2, F2	IF	ID	EX	ME WB
DIVD	F3, F3, F30	(	MUCHOS		
SD	F3, O(R3)	STALUS'	sep or		



# 1) Superescalar VS VLIW

Superescalares pueden ir en desorden. VLIW se llevan a cabo en orden, salvo end port, ellos llevan su orden

LD F2,0(R1)	• Fully pipelined		
MULD F4,F2,FO	• Branch delay ignored		
LD F6,0(R2)			
ADDD F6,F4,F6	ALU	ALU	3
SD F6,0(R2)	ALU	SD	2
DADDUI R1,R1,-8	LD	ALU	1
DADPUI R2,R2,-8	SD	LD	
SLT R3,R1,end	Fixed $\rightarrow$ FO		
BEQZ R3,loop	Unroll 4 times		

VLIW  
(order)

LD F2,0(R1)	LD F8,-8(R1)	LD F14,-16(R1)	LD F20,-24(R1)
MULD F4,F2,FO	MULD F10,F8,FO	MULD F16,F14,FO	MULD F22,F20,FO
LD F6,0(R2)	LD F12,-8(R2)	LD F18,-16(R2)	LD F24,-24(R2)
ADDD F6,F4,F6	ADD F12,F10,F12	ADD F18,F16,F18	ADD F24,F22,F24
SD F6,0(R2)	SD F12,-8(R2)	SD F18,-16(R2)	SD F24,-24(R4)
	DADDUI R1,R1,-32		
	DADPUI R2,R2,-32		
	SLT R3,R1,end		
	BEQZ R3,loop		

cycle	MEM 1	MEM 2	FP1	FP2	Int/branch
1	LD F2,0(R1)	LD F6,0(R2)			
2	LDF8,-8(R1)	LD F12,-8(R2)	NOP LD $\rightarrow$ ALU		
3	LDF14,-16(R1)	LD F18,-16(R2)	MULD F4,F2,FO	MULD F10,F8,FO	
4	LD F20,-24(R1)	LD F24,-24(R2)	MULD F16,F14,FO	MULD F22,F20,FO	
5			NOP ALU $\rightarrow$ ALU		DADDUI R1,R1,-32 (*)
6			" "		DADDUI R2,R2,-32
7			ADD F6,F4,F6	ADD F12,F10,F12	
8			NOP ALU $\rightarrow$ SD	ADD F18,F16,F18	
9			NOP ALU $\rightarrow$ SD	ADD F24,F22,F24	
10	SD F6,0(R2) <sup>32</sup>	SD F12,-8(R2)			SLT R3,R1,end
11	SD F18,-16(R2) <sup>16</sup>	SD F24,-24(R2) <sup>8</sup>			BNE R3,loop

(\*) Al haber restado 32, se debe compensar, sumando 32 a los SD

## 2) Poison bits

LW\* R1, 0(R2)  
if-1: BEQZ R3, endif-1

R1	R2	R3	R4	RS	R6	R7	R8	R9
10	0	0	4	5	6	7	8	9

DADDI R1, R4, 8

Access to MEM position 0 is forbidden, no LD or SD

endif-1: LW\* RS, 0(R2)

if-2: BEQZ R6, endif-2

DADDI RS, R7, -1

endif-2: DADD R8, R1, RS

SW R8, 0(R2)

DADDI R9, R0, 1

a) Processor with no speculation help

Instruc	Produces Exception? Y/N	Terminates?	Performs
LW* R1, 0(R2)	Y	Y	R1 ← <del>Undef</del>

b) HW/SW cooperation. Not distinguish between norm and spec. instrcts. Always returns undef.

Instruc	Produces Exception? Y/N	Terminates?	Performs
LW* R1, 0(R2)	Y	N	R1 ← Undef
BEQZ R3, endif-1	N	N	Jumps
LW* RS, 0(R2)	Y	N	RS ← Undef
BEQZ R6, endif-2	N	N	Doesn't jump
DADDI RS, R7, -1	N	N	RS ← 7 - 1 = 6
DADD R8, R1, RS	N	N	R8 ← Undef
SW R8, 0(R2)	Y	Y	No storage
DADDI R9, R0, 1	N	N	R9 ← 10 + 1

c) HW/SW cooperation. Distinguish between norm and spec. Normal terminates

Instruc	Produces Exception? Y/N	Terminates?	Performs
LW* R2, 0(R2)	Y	N	R2 ← Undef
BEQZ R3, endif-1	N	N	Jumps
LW* RS, 0(R2)	Y	N	RS ← Undef
BEQZ R6, endif-2	N	N	Doesn't jump
DADDI RS, R7, -1	N	N	RS ← 7 - 1
DADD R8, R1, RS	N	Y	R8 ← Undef
SW R8, 0(R2)	Y	Y	No store

Instruc	Produces Exception? Y/N	Terminates?	Performs
LW* R1, 0(R2)	Y	N	R1 ← dirty
BEQZ R3, endif-1	N	N	Jumps
LW* RS, 0(R2)	Y	N	RS ← dirty
BEQZ R6, endif-2	N	N	Doesn't jump
DADDI RS, R7, -1	N	Y	RS ← 7 - 1 = 6
DADD R8, R1, RS	Y	Y	Terminates

R1	R2	R3	R4	RS	R6	R7	R8	R9
Dirty	0	0	4	6	6	7	8	9

- a) Mark the stalls that exist in the code from I6 to I13 (in a single iteration). Indicate between which instructions they occur and the reason why they occur.

Name: \_\_\_\_\_

**Important:** Solve each exercise on a different sheet. Do not forget to write down your name on all the sheets you deliver.

**Exercise 1 (2.5 points)**: The following MIPS64 code calculates the dot product of two vectors (vector1 and vector2). The dot product result is stored in \$T0 at the end of the loop.

```
;Example vectors
vector1:    .float 1.0, 2.0, 3.0, 4.0, 5.0, 6.0
vector2:    .float 6.0, 5.0, 4.0, 3.0, 2.0, 1.0

;Variable initialization
addi $t0, $zero, 0      ; I1 (i = 0)
addi $t1, $zero, 6       ; I2 (n = 6)
add.s $f0, $zero, $f31   ; I3 ($T0=0)
la $t2, vector1         ; I4 (get address of vector1)
la $t3, vector2         ; I5 (get address of vector2)

Loop:
    l.s $f4, ($t2)        ; I6 (read element from vector1)
    l.s $f6, ($t3)        ; I7 (read element from vector2)
    mul.s $f8, $f4, $f6     ; I8 (multiply the values)
    add.s $f0, $f0, $f8     ; I9 (add the result to the dot product)
    addi $t0, $t0, 1          ; I10 (i = i + 1)
    addi $t2, $t2, 4          ; I11 (get next position of vector1)
    addi $t3, $t3, 4          ; I12 (get next position of vector2)
    blt $t0, $t1, Loop      ; I13 (if i < n, go back to the beginning)
```

The computer used has the following characteristics:

- A data value can be written to and read from a register in the same clock cycle.
- Operand forwarding is used. **(\*)**
- All dependencies are detected and resolved in the ID stage.
- Branches are entirely resolved in the ID stage. Delayed branching is applied.
- Two WB operations in floating point are not allowed in the same cycle.
- In the same cycle, one floating point WB operation and another integer operation can be performed.
- All functional units are segmented.

Here is the table that shows the number of cycles used by each functional unit:

Functional Unit	Number of cycles used
FP adder/subtractor	2
FP multiplier	4
FP Divider	12
Other units	1

- b) Calculate the CPI of the program (consider all the instructions executed from I1; show the operations performed).

- c) Rename and unroll the loop the minimum number of times to schedule it without stalls. The resulting code should perform the same operation as the original code. Then, answer the following questions:

a. What is the CPI of the resulting code? (show the operations performed)

b. Complete the following table with the resulting code.

Between which instructions	Reason
I6 → I8	Data Dependency through f4
I7 → I8	Data Dependency through f6
I8 → I9	

$$\begin{aligned} \text{CPI} &= \frac{\text{Total cycles}}{\text{Total instruc}} = \frac{5+12 \cdot 3}{5+12 \cdot 3} \\ &= \underline{\underline{1}} \end{aligned}$$

**Important:** Solve each exercise on a different sheet. Do not forget to write down your name on all the sheets you deliver.

**Exercise 2 (2,5 points).** Let be a MIPS64 processor using hardware speculation and dynamic scheduling using the Tomasulo technique with speculation, with the following characteristics:

- The load and store execution stage takes 2 cycles.
- The floating-point (FP) multiplication execution step takes 7 cycles.
- The floating-point addition (FP) execution step takes 3 cycles.
- A 4-entry reorder buffer (ROB) is provided.
- 2-entry reservation stations are available for each unit in FP.
- It has different load and store buffers of 2 entries each.
- Any stage of an instruction that wants to write to the CDB (Common Data Bus) cannot do so if another instruction is already using it.
- In general, any resource released at a given clock cycle cannot be used until the next one.

- a) Detail the execution of the sequence of instructions shown below when using the Tomasulo algorithm with speculation. To do so, indicate in the table below the cycle number (start and end if applicable) for each stage of each instruction.

Instruction	Issue	Execution	Results Writing	Commit
L.D F0, 0(R1)	1	2-3	4	5
L.D F2, 0(R2)	2	3-4	5	6
ADD D F4, F0, F2	3	6-7	8	9
L.D F6, 0(R3)	5	10-11	12	13
MUL D F8, F4, F6	6	13-14	15	16
ADD D F10, F6, F12	7	16-17	18	19
S.D F8, 8(R1)	11	19-20	21	22
S.D F10, 8(R2)	12	20-21	22	23

- b) Fill the table below to show the contents of the reorder buffer (ROB) at cycle 10. Use only the necessary table entries.

→ "Aquellos en ejecución o WB en el ciclo 10"

Entry #	Instruction	Stage	Destination Register	Value
1	ADD F4, F0, F2	Commit	F4	F0 + F2
2	L.D F6, 0(R3)	Wait commit	F6	MEM[0R3]
3	MUL D F8, F4, F6	Exec	F8	F4 * F6
4	ADD D F10, F6, F12	Exec	F10	F6 + F12

**Important:** Solve each exercise on a different sheet. Do not forget to write down your name on all the sheets you deliver.

**Exercise 3 (2 points)** Consider a MIPS64-based processor, and various mechanisms for handling exceptions produced by speculative instructions. From the following code:

```

DADDI R1, R1, #1
BEQ R1, R2, then
        DDIV R5, R1, R5
        end:
else:
        DDIV R5, R2, R1
        DADDI R3, R3, #1
then:
        DADD R4, R4, R3
end:

```

Considering that there is a high probability that the branch condition will not be met, the compiler has generated the next version of the same code, with one speculative instruction (highlighted):

- Indicate, depending on the exception handling mechanism available in the processor, and for the speculative version of the code, the instructions that cause any exception, and whether or not they terminate the program, as well as the value of the R1-R5 registers at the end of the program, knowing that their initial state is:
- |    |    |    |    |    |
|----|----|----|----|----|
| R1 | R2 | R3 | R4 | R5 |
| 0  | 1  | 2  | 3  | 0  |
- The processor has no speculation support, i.e., it does not distinguish between normal and speculative instructions. Therefore, any instruction that throws an exception will cause the program to be terminated immediately.
  - The processor has a hardware/software cooperation mechanism that distinguishes between normal and speculative instructions. If the instruction that throws the exception is normal, the exception's service routine will terminate the program, but if it is speculative, the routine will return an undefined value (without aborting the execution).
  - Speculation with poison bits.

Name: Javier García Tercero

Indicate IN THE NEXT TABLE the correct answer to each test question below:

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
A	B	A	B	C	D	A	B	C	A	B	C	A	C	A
16	17	18	19	20	21	22	23	24	25	26	27	28	29	30
C	C	D	A	C	C	A	C	B	D	A	A	D	B	

**Important:** Test counts 3 points. Each correct question counts 0.1 points. Wrong answers subtract 0.1/3 points. Not answered questions are not considered for the mark.

- Which of the following statements is correct?
  - An organization can survive for many years, but its architecture changes as technology evolves
  - Many computer manufacturers offer a family of models, all with the same organization, but with a difference in architecture
  - The architecture of a computer must be designed to implement a specification of a particular organization
  - All statements are false
- The Linpack benchmark used in the TOP500 is of the type:
  - Synthetic
  - Real program
  - Kernels
  - Suites
- From compiler X we have obtained an improved compiler M, which provides a code for floating-point instructions that is k times faster than the original. If 40% of the total time is spent on floating-point instructions when using the C compiler, what would be the maximum speed increase we could obtain with compiler M?
  - 1.67
  - 167%
  - 2
  - All are false
- We want to analyse the performance of two computers with the same instruction set. In that set, there are three types of instructions: A, B and C. Computer 1 needs 2, 4 and 6 cycles to execute instructions of type A, B and C, respectively. Computer 2 needs 2, 2 and 5 cycles. Computer 1 operates at 1.2 GHz and computer 2 at 1.8 GHz. The benchmark used is composed of 20% of type A instructions, 50% of type B and 30% of type C. What is the MIPS speed of both computers?
  - Computer 1: 285.7 and Computer 2: 512.2
  - Computer 1: 285.7 and Computer 2: 620.7
  - Computer 1: 222.2 and Computer 2: 610.2
  - All are false
- What problem does "register renaming" solve?
  - Control dependencies
  - Cache misses
  - Name dependencies
  - All are correct
- In a basic segmented processor with forwarding, we execute this code:
 

```
Loop:    L.D F8, 0(R1) ;I1
              ADD.D F4, F2, F8 ;I2
              SUB.D F4, F2, F0 ;I3
              ADD.D F4, F4, F0 ;I3
```

Knowing that the addition and subtraction functional units take 2 cycles. Indicate the correct option:

  - There are two output dependencies
  - There are two anti-dependencies
  - There are two data dependencies
  - All are correct
- Consider the following loop:
 

```
Loop:    L.D F8, 0(R1) ;I1
              ADD.D F6, F2, F8 ;I2
              S.D F6, 0(R1) ;I3
              DADDUI R1, R1, #-8 ;I4
              BNEZ R1, Loop ;I5
```

Between the first and second instruction (L.D and ADD.D):
 
  - There is a data dependency in the same iteration
  - There is an output dependency in the same iteration
  - There is a data dependency between iterations
  - None are correct
- In a new segmented processor of 10 stages with forwarding, all possible hardware improvements have been applied to avoid stalls. The CPI (Cycles Per Instruction) ideally achievable for the system is:
  - 0.1
  - 1
  - 10
  - All are false
- The Scoreboard technique is used to...
  - Dynamically rename registers
  - Dynamically resolve data dependencies
  - Reduce the number of instructions to execute
  - All are false
- What is the technique called that allows instructions to execute out of order, but ensures that the results are generated according to the program order?
  - Speculative programming
  - Compiler programming
  - Dynamic programming
  - Static programming
- What is a reservation station in the context of the Tomasulo algorithm?
  - A physical storage location for data that allows their speculation
  - A physical storage location for registers that allows their renaming
  - A physical storage location for data dependencies that allows their resolution
  - All are correct
- Given a 5-stage pipeline in which 2 out of every 9 instructions is a taken branch, and 2 delay slots happen for each taken branch, indicate which of the following is correct:
  - CPI = 1.44
  - CPI = 1.22
  - CPI = 0.81
  - CPI = 0.6
  - CPI = 1

$$CPI = \frac{\text{Total of delay}}{\text{Instructions}} = \frac{9+4}{9} = 1.44$$

Addition cycles:  $2 \times 2 = 4$

↳ delay per instruction

$$CPI = \frac{1}{\text{IPC}} = \frac{1}{0.69} = 1.44$$

14. Indicate which of the following statements is correct:
- If the branch condition is evaluated in the EX stage and the branch target address is computed in the MEM stage, we will have a scheme with 1 delay slot
  - If the branch condition and address jump are both computed in the ID stage, we will have a scheme with 3 delay slots
  - If the branch condition is evaluated in the EX stage and the branch target address is computed in the ID stage, we will have a scheme with 2 delay slots
  - All of the above are correct

23. Each iteration of a given loop in a scalar processor requires 15 clock cycles to be issued. We unrolled such a loop 5 times and scheduled it for a 4-way superscalar obtaining that we now perform 5 iterations of the original loop in 15 cycles. What is the speedup obtained?
- $$\text{Speedup} = \frac{\text{Time before}}{\text{Time after}} = \frac{15 \times 5}{15} = 5$$

15. What is the latency of a branch instruction?
- The number of additional cycles between the IF stage and the cycle of its complete resolution
  - The number of cycles between the IF stage and the ID stage
  - The number of additional cycles between the IF stage and the cycle where its condition is computed
  - The number of cycles between the IF stage and the MEM stage
16. Dynamic branch prediction techniques focus on:
- the behavior through the statistics of all the branch instructions
  - the use of software compilers that predict the behavior of a program
  - speculating through poison bits which branch instructions are to be taken
  - the behavior of the last executions of branch instructions
17. A "speculative" instruction:
- Is a conditional branch whose condition must be predicted
  - Is an instruction that causes a non-recoverable exception
  - It is an instruction executed before knowing whether it should be executed
  - An instruction that calculates a probable result of an arithmetic-logical operation
18. When a branch instruction that has been mispredicted reaches the top of the ROB:
- The data structures used for the prediction are updated
  - The fetch of the correct instruction is initiated
  - The ROB is flushed
  - All are correct
19. The problem of memory "disambiguation":
- Refers to the occurrence of potential RAW hazards in memory
  - Refers to the occurrence of potential WAR hazards in memory
  - Refers to the occurrence of potential in-memory WAW hazards
  - Refers to the occurrence of potential structural hazards in memory
20. We need to schedule a sequential code consisting of 40 instructions for a VLIW processor with 4 operations per instruction. In the best case scenario (assuming an "ideal" code and processor), we would have:
- 40 long instructions
  - 20 long instructions
  - 10 long instructions
  - 4 long instructions
21. A characteristic of superscalar processors is...
- They use long instructions and dynamic scheduling
  - They always employ static prediction
  - They execute multiple instructions per cycle
  - Instructions are always retired out-of-order
22. In a superscalar processor, in the event of an exception in the *Execution* stage...
- The Reorder Buffer (ROB) is immediately emptied, preventing the execution of more program instructions
  - The resources (ROB and others) used by the instruction that caused the exception are released, but the program execution continues normally
  - It immediately jumps to the exception handler
  - None of the above

Use the tables below to present the results of this part. Use only the rows you need in each case. In a) and b) use the label “undefined” to indicate that the value of a register is undefined, and in c) use the label “dirty” to indicate that the *poison bit* associated to a register is set (equal to 1).

a) Processor without speculation support.

Instruction executed	$\{\text{Exception raised?}\}$ (Yes/No)	$\{\text{Program termination?}\}$ (Yes/No)	Result (jump/don't jump, dest. reg. value...)	Result (jump/don't jump, dest. reg. value...)
DADDI R1, R1, #1	No	No	R1 $\leftarrow$ 1	R1 $\leftarrow$ 1
DDIV* R5, R1, R5	Yes	Yes	Program terminates	Program terminates

b) Processor without speculation support.

Register file content after program termination:				
R1	R2	R3	R4	R5
1	1	2	3	0

b) Hardware/software cooperation mechanism.

Instruction executed	$\{\text{Exception raised?}\}$ (Yes/No)	$\{\text{Program termination?}\}$ (Yes/No)	Result (jump/don't jump, dest. reg. value...)	Result (jump/don't jump, dest. reg. value...)
DADDI R1, R1, #1	No	No	R1 $\leftarrow$ 1	R1 $\leftarrow$ 1
DDIV* R5, R1, R5	Yes	No	R5 $\leftarrow$ undefined	R5 $\leftarrow$ undefined
BNE R1, R2, end	No	No	Jumps	Jumps
DDIV R5, R2, R1	No	No	R4 $\leftarrow$ 5, program terminates	R4 $\leftarrow$ 5, program terminates
DADD R3, R3, #1	No	No		
DADD R4, R4, R3	No	Yes	R4 $\leftarrow$ 6, program terminates	R4 $\leftarrow$ 6, program terminates

b) Hardware/software cooperation mechanism.

Register file content after program termination:				
R1	R2	R3	R4	R5
1	0	2	3	0

c) Speculation with *poison bits*.

Instruction executed	$\{\text{Exception raised?}\}$ (Yes/No)	$\{\text{Program termination?}\}$ (Yes/No)	Result (jump/don't jump, dest. reg. value...)	Result (jump/don't jump, dest. reg. value...)
DADDI R1, R1, #1	No	No	R1 $\leftarrow$ 1	R1 $\leftarrow$ 1
DDIV* R5, R1, R5	Yes	Yes	R5 $\leftarrow$ dirty (poison bit=1)	R5 $\leftarrow$ dirty (poison bit=1)
BNE R1, R2, end	No	No	Jumps	Jumps
DADD R4, R4, R3	No	Yes	R4 $\leftarrow$ 5, program terminates	R4 $\leftarrow$ 5, program terminates

c) Speculation with *poison bits*.

Instruction executed	$\{\text{Exception raised?}\}$ (Yes/No)	$\{\text{Program termination?}\}$ (Yes/No)	Result (jump/don't jump, dest. reg. value...)	Result (jump/don't jump, dest. reg. value...)
DADDI R1, R1, #1	No	No	R1 $\leftarrow$ 1	R1 $\leftarrow$ 1
DDIV* R5, R1, R5	Yes	No	R5 $\leftarrow$ dirty (poison bit=1)	R5 $\leftarrow$ dirty (poison bit=1)
BNE R1, R2, end	No	No	Doesn't jump	Doesn't jump
DDIV R3, R3, #1	No	No	R4 $\leftarrow$ 3 (poison bit=0)	R4 $\leftarrow$ 3 (poison bit=0)
DADD R4, R4, R3	No	Yes	R4 $\leftarrow$ 6, program terminates	R4 $\leftarrow$ 6, program terminates

c) Speculation with *poison bits*.

Register file content after program termination:				
R1	R2	R3	R4	R5
1	0	2	5	dirty

Speculation part:

In part 3 of the previous part (use the tables below), but now starting from these initial values for registers R1-R5:



Jan 24

Exer 3.

DADDI R1, R1, #1  
 DDIV\* R5, R1, RS  
 BNE R1, R2, end  
 DDIV R5, R2, R1  
 DADDI R3, R3, #1  
 end: DADD R4, R4, R3

; speculative division

1st Part	R1	R2	R3	R4	RS
	0	1	2	3	0

a) No spec. Support. If except = terminate

Instruc exec.

Except Raised? Y/N

Program Term? Y/N

Result

R1 ← 1

Program terminates due to 0/0

DADDI R1, R1, #1

N

N

DDIV\* R5, R1, RS (R1/RS)

Y

Y

b) HW+SW cooperation, distinguishes between normal and speculative instruc, if except normal → terminates if specul. → undef value

Instruc exec

Except Raised? Y/N

Program Term? Y/N

Result

R1 ← 1

RS ← undef

No jmp; R1 = R2 (1:1)

R5 ← 1; R2/R1 (1/1)

R3 ← 3; 2 + 1

R4 ← 6; 3 + 3; end program

DADDI R1, R1, #1

N

N

DDIV\* R5, R1, RS

Y

N

BNE R1, R2, end

N

N

DDIV R5, R2, R1

N

N

DADDI R3, R3, #1

N

N

DADD R4, R4, R3

N

Y

c) Speculative with poison bits

Instruc

Except Raised? Y/N

Program Term? Y/N

Result

R1 ← 1

RS ← dirty (poison bit: 1)

R1 = R2; doesn't jump

RS ← 1 (poison bit: 0)

R3 ← 3

R4 ← 6

DADDI R1, R1, #1

N

N

DDIV\* R5, R1, RS

Y

N

BNE R1, R2, end

N

N

DDIV R5, R2, R1

N

N

DADDI R3, R3, #1

N

N

DADD R4, R4, R3

N

Y

NEW R1 R2 R3 R4 RS  
 0 0 2 3 0

b) HW+SW cooperation

Instruc Exec

Except Raised? Y/N

Program Term? Y/N

Result

R1 ← 1

RS ← undef

{ R1 1

  R2 0

  R3 2

  R4 5

  RS undef

DADDI R1, R1, #1

N

N

DDIV\* R5, R1, RS

Y

N

BNE R1, R2, end

N

N

DDIV R5, R2, R1

N

N

DADD R4, R4, R3

N

Y

c) speculative with poison bits

DADDI R1, R1, #1

N

N

DDIV R5, R1, RS

Y

N

BNE R1, R2, end

N

N

DADD R4, R4, R3

N

Y

R1 ← 1

RS ← dirty (poison bit: 1)

jumps

RS ← 5



# Branch Based on the exec history

## 1-Bit predict

Bit stores behavior of last exec.

N T (0)  
T (1)

## 2-Bit predict

If predict 2 times wrong, changes

SNT (00)  
WNT (01)  
WT (10)  
ST (11)

normally, predictor n=2 is equal or better than n>2

## (m, n) predictor

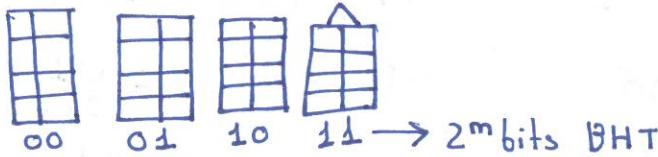
Uses info from m branches

Used among  $2^m$  branch predictions, being each a n-bit predictor for a single branch

BHR : Branch History Register  $\rightarrow 2^m$  BHT tables

BHT : Branch History Table  $\rightarrow$  n bits per entry      n bits

Example: (2,2) predictor



## Exer 8

Hybrid. 2 bit local predictor. Entry per branch. Global (1,2), 2 entry BHT entries indexed by the least significant bit of branch address. Selector 2-bit saturating ( $\uparrow$  Taken,  $\downarrow$  NT). Local when selector = 0,1, global = 2,3

loop: LW R2, 0(R1) 6

ANDI R4, R2, #1 0

BEQZ R4, next<sub>T</sub> (B1, branch B1 = 0x654)

DADDUI R1, R1, #4 9

DADDUI R2, R2, #1

R1 and R3 initialized with: 3, 2, 5, 4, 6

11	10	0101	0100	0110
01	01	0001	0001	0002
01	00	0001	0000	0000

next: BNE R1, R3, loop (B2, branch B2, = 0x623)

SW R2, -4(R1) 5

WNT

Branch	ACTUAL Predictor	Selector	Local Predictor	Global Predictor
Initial state	-	00	01 01	BHR=0      81 81      81
B1(v=3)	NT	NT(h) 00	00 00	BHR=0      81 81      81
B2(v=3)	T	NT(m) 01	01 01	BHR=0      01 01      01
B1(v=2)	T	NT(m) 10	01 10	BHR=0      01 01      01
B2(v=2)	T	NT(m) 11	01 10	BHR=-1      01 10      01
B1(v=5)	NT	NT(h) 10	02 10	BHR=0      02 10      01
32(v=5)	T	T(h) 11	01 10	BHR=1      01 11      00
B1(v=4)	T	NT(m) 11	01 10	BHR=1      01 11      02
B2(v=4)	T	NT(m) 11	01 10	BHR=4      01 11      10
B1(v=6)	T	NT(m) 11 (Miss aqui GP)	01 10	BHR=1 (*) 01      10 11      10
B2(v=6)	NT	T(m) 10	01 10	BHR=0      01      10 11      01

(\*) Miss en la tabla 1, B1 defence. Establece tabla 2.

## 6 First iterations

DADD R1, R0, R0

L1: BEQZ R1, L2

...

#B1

L2: BEQZ R1, L3

...

#B2

.3: XORI R1, R1, #1

#(↔NOT RL)

J L1

$$\begin{array}{r} 0101 \\ 0001 \\ \hline 0100 \end{array}$$

$$\begin{array}{r} 0000 \\ 0001 \\ \hline 0001 \end{array}$$

"Esto va ↓  
en pasado"

2<sup>m</sup> tables = 4

&lt; 00 &gt;

01

11

10

(2,2)

n bits

a) Actual behaviour  
B1/B2      Actual  
initial state   T/NTb) One-level predict. behavior  
Predict bits for B1, B2      Predicted T/NT      Hit Y/Nc) Two-level predictor behavior  
Prediction bits (BHR and BTHs)      Predicted outcome (T/NT)      Hit (Y/N)

31(R1=0) T

-



$$\begin{array}{ccccccccc} & & & & & & & & \\ & 00 & 00 & 00 & 00 & 00 & - & - & \\ & 00 & 00 & 00 & 00 & 00 & & & \\ & 01 & 01 & 00 & 00 & 00 & & & \\ & 00 & 00 & 00 & 00 & 00 & & & \\ & 00 & 00 & 00 & 00 & 00 & & & \end{array}$$

NT N(m)

NT N(m)

NT Y(h)

NT Y(h)

NT N(m)

NT N(m)

NT Y(h)

NT Y(h)

NT N(m)

NT N(m)

NT N(m)

NT N(m)

NT N(m)

NT N(m)

11	10	00	00	00
00	10	00	00	00
10	10	00	00	00
00	10	00	00	00

Por el BHR anterior apuntamos a la tabla 3.  
Restariamos 1 por ser NT el actual pero como es  
ya, ahí se queda

El BHR sirve para saber qué tabla activaríamos  
en el momento. Se suma 1 si el actual fue taken.  
Se resta 1 si fue not taken

Nov 2015

- 2 Bit local predictor, entry per branch instruc  $\rightarrow 2, B1 \text{ y } B2$
- Global (2,1) predictor, each BHT table contains one entry per branch instruc
- Selector of 2-bit saturating counter. Taken $\uparrow$ ; Not Taken $\downarrow$ . Local when selector = 0 or 1, global = 2.
- Selector initialized to 11, rest = 0

R1 and R3 initialized to the sequence: 0, 0, 1, 0, 0, 2

loop: LW R2, 0(R1)

BEQZ R2, skip #B1

SW R0, 0(R1)

DADDI R4, R4, #1

skip: BEQ R1, R3, end #B2

DADDI R1, R1, #4

j loop

end: ...

Branch exec. (B1 or B2)	Actual T/NT	Predictor T/NT(h/m)	Selector	Local Predictor	Global Predictor
B1 (R2=0)	T	NT(m)	11	0000	BHR=00 BHTs: 00 00 00 00 BHR: 01 permss ↙

Step	B1 or B2	R1	R4	T/NT	Predictor	HIT	BHR	Table 1		Table 2		Table 3		Table 4		Predict	HIT
								01	00	00	00	00	00	00	00		
1	B1	0	?	NT	01	Y	00	00	00	00	00	00	00	00	00	NT	Y
2	B2	0	0	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
3	B2	0	1	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
4	B2	0	2	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
5	B2	0	3	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
6	B2	0	4	T	00	N	01	01	00	00	00	00	00	00	00	NT	N
7	B1	1	4	NT	00	Y	10	01	00	00	00	00	00	00	00	NT	Y
8	B2	1	0	NT	00	Y	00	01	00	00	00	00	00	00	00	NT	Y
9	B2	1	1	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
10	B2	1	2	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
11	B2	1	3	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
12	B2	1	4	T	00	N	01	01	00	00	00	00	00	00	00	NT	N
13	B1	2	4	NT	01	Y	10	01	00	00	00	00	00	00	00	NT	Y
14	B2	2	0	NT	00	Y	00	01	00	00	00	00	00	00	00	NT	Y
15	B2	2	1	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
16	B2	2	2	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
17	B2	2	3	NT	00	Y	00	00	00	00	00	00	00	00	00	NT	Y
18	B2	2	4	T	00	N	01	01	00	00	00	00	00	00	00	NT	N
19	B1	3	4	T	01	N	11	00	01	00	00	00	00	00	00	NT	N

hit rate:  $\frac{15}{19} \cdot 100 = 78.94\%$

hit rate:  $\frac{15}{19} \cdot 100 = 78.94\%$

$$\frac{15 \text{ hits}}{19 \text{ hits + miss}} = 0.78 \cdot \frac{1}{miss} = 0.78 \text{ hits}$$

BHR = Branch History Register  
(1,2)

sbt rd, rs, rt  
if rs < rt  
  rd = 1  
if rs > rt  
  rd = 0

Loop: LS F2, O(R2)  
 MULS F4, F2, F2  
 MULS F6, F4, F1  
 S.S F6, O(R4)  
 DADDUI R2, R2, 4  
 DADDUI R4, R4, 4  
 DADDUI R3, R3, 1  
 BNE R3, R1, loop

	ALU	ALU	Int	LD	LD
1					
2					
3					
4					
5					
6					
7					
8					
9					
10					

Latencia → Huecos entre instruc. tal cual  
 3 ↳ Si fuese Tiempo de ejecución, sería los huecos + la propia instrucción contada  
 2  
 0  
 1  
 0  
 • No delay ↳ No hay que poner una instrucción debajo del BNE

ciclo	MEM	MEM	FP	FP	Int / branch
1	LS F2, O(R2)				
2					DADDUI R2, R2, 4
3					DADUI R4, R4, 4
4		MULS F4, F2, F2			DADDUI R3, R3, 1
5					
6					
7		MULS F6, F4, F1			
8					
9					
10	S.S F6, -4(R4)				BNE R3, R1, loop

Instruc	Issue	Execution	Write	Commit
MULD F0, F4, F12	1	2-5	6	7
SUBD F4, F0, F10	2	7-8	9	10
LD F0, O(R1)	3	4	5	11
DIVD F0, F8, F14	4	6-25	26	27
SUBD F12, F8, F0	8	27-28	29	30
ADDD F10, F14, F24	11	12-13	14	31
LD F0, O(R2)	12	13	15	32
MULD F18, F14, F24	28	29-32	33	34

4 ROB

2+ 4 x 20 %

1 FU pa cada uno

CDB+1 mínimo

1<sup>er</sup> LD (1)

ROB en ciclo que finaliza

1<sup>er</sup> LD (1)

- 0 DIVD F0 F8 F14
- 1 SUBD F12 F8 F0
- 2 LD F0 O(R3)
- 3 ADD F10 F14 F24

Exec  
Wait issue  
Commit  
Issue

F8/F14  
F8-F0  
MEM F0 → R1  
F14+F24

Local 2 bits global of (1,2). Selector 2 bit saturated. T↑ NT↓  
Local when 0 or 1 global 2 or 3. SOLO ACTUALIZA el predictor empleado

loop: LW R2, 0(R1)  
ANDI R4, R2, 1  
BEQZ R4, next

DADDUI R1 R1 4

DADDUI R2 R2 1

next: BNE R1, R3, loop  
SW R2, 4(R1)

3, 2, 5, 4, 6

011	010
001	001
-----	
001	000

B1

→ Cond 1 ya no avanza  
B2 más, habrá llegado  
al final y nos dará

2<sup>m</sup> n bits

2<sup>1</sup> 2 bits

weak NT

Branch	T / NT?	Predicts	Local predictor	Selector	Global BHR	Predictor BHT0	Predictor BHT1
-	-	-	01 01	00	0	01 01	01 01
B1(3)	NT	NT	00 01	00	0	01 01	01 01
B2(3)	T	NT	00 01	01	0	01 01	01 01
B1(2)	T	NT	01 10	10	0	01 01	01 01
B2(2)	T	NT	01 10	11	1	01 01	01 01
B1(5)	NT	NT	01 10	10	0	01 10	00 01
B2(5)	T	T	01 10	11	1	01 11	00 01
B1(4)	T	NT	01 10	11	1	01 11	01 01
B2(4)	T	NT	01 10	11	1	01 11	01 10
B1(6)	T	NT	01 10	11	1	01 11	10 10
B2(6)	NT	T	01 10	10	0	01 11	10 01

Fórmulas

$$\text{Speedup} = \frac{1}{(1-f) + f/s}$$

$$\text{MIPS} = \frac{\text{clock freq}}{\text{CPI} \cdot 10^6}$$

$$\text{speedup} = \frac{\text{before}}{\text{after}}$$

$$\text{CPI} = \frac{\text{total cycles}}{\text{total instruc}}$$

$$\text{IPC} = 1/\text{CPI}$$

Flynn Taxonomy

Uniproc ← SISD

MISD → No use TPU

SIMD  
Vector, multimedia  
GPU

MIMD  
→ SMP, NUMA, cluster, cloud

**SCHOOL OF COMPUTER SCIENCE AND ENGINEERING**  
**COMPUTING SYSTEMS DEPARTMENT**  
**COMPUTER SCIENCE DEGREE. COMPUTER ARCHITECTURE**

**Final written exam (June 2023)**      **Test**      **Year 2022/23**



Name: \_\_\_\_\_

Indicate the correct answer to each test question **IN THE NEXT TABLE**:

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23	24	25	26	27	28	29	30

**Important:** Test counts **3 points**. Each correct question counts 0.1 points. Wrong answers subtract 0.1/3 points. Not answered questions are not considered for the mark.

1. The peak performance of an architecture is:
    - a. The performance of a machine is guaranteed not to exceed
    - b. The performance of a machine when running the Linpack benchmark
    - c. The maximum performance a machine achieves with the most favourable benchmark
    - d. All of them are false
  2. The CPI depends on:
    - a. The architecture of the instruction set
    - b. The CPU organisation
    - c. On the complexity of the set of instructions
    - d. All of them are true
  3. Suppose you have two machine realisations of the same instruction set architecture. Machine A has a clock cycle of 1 ns and a CPI of 3 for a particular program, while machine B has a cycle time of 2 ns and a CPI of 1.5 for the same program. Which machine is faster for this program, how much faster?
    - a. A is 20% faster than B
    - b. B is 20% faster than A
    - c. A is as fast as B
    - d. All of them are false
  4. We want to analyse the performance of a machine. In its instruction set there are three types of instructions, A, B and C, and it needs 2, 6, and 4 clock cycles to execute instructions of type A, B, and C. The machine has a clock frequency of 800 MHz. The test program we are going to use has the following mix of instructions: 40% type A, 50% type B, and 10% type C. What is the peak MIPS speed of the machine?
    - a. 190,5
    - b. 200
    - c. 400
    - d. None of the above
  5. Renaming floating point records allows to avoid:
    - a. Data dependencies.
    - b. Name dependencies.
    - c. Control Dependencies.
    - d. None of the above.
  6. Mark which of the following statements is correct for MIPS64:
    - a. Pipelining decreases the latency of a task.
    - b. It uses one cache memory for instructions and another for data.
    - c. The PC is updated at the final stage of each instruction (WB).
    - d. None of the above.
  7. Select the correct option for the following loop and the ADD/D and SD instructions:
 

```
loop: L.D F1, 0(R1)
          ADD.D F2, F1, F0
          BNEZ R1, Loop
```
8. The following code shows the modified loop that resulted from applying the software pipelining technique to an initially given code.
- ```
Loop: S.D F4, 16(R1) ; 1 store M[1]
        ADD.D F4, F0, F2 ; 2 add a M[i-1]
        L.D F0, 0(R1) ; 3 load M[i-2]
        DADDI R1, R1, #8 ; 4 update counter
        BNEZ R1, Loop ; 5 jump if R1 = 0
```
9. In the Scoreboard technique, the results of the functional units:
- a. They are sent to the instructions that are waiting for them through the CDR.
  - b. The initial code has the DADDUI instruction with an increment of #8.
  - c. The initial code has the DADDUI instruction with an increment of #4.
  - d. The initial code does not have the DADDUI instruction.
10. Mark the correct option for dynamic instruction scheduling without speculation:
- a. Issue in order, execution out of order and completion in order.
  - b. Issue in order, execution out of order and completion out of order.
  - c. Issue in order, execution out of order and completion out of order.
  - d. All are false.
11. Given a 5-stage pipeline in which 2 out of every 8 instructions are taken branches and there are 3 delay slots for each taken branch, indicate which of the following options is correct:
- a. CPI = 1.2, IPC = 0.9
  - b. CPI = 1.3, IPC = 0.75
  - c. CPI = 1.75, IPC = 0.57
  - d. CPI = 1.1, IPC = 0.95
12. A code is executed on a 5-stage pipelined processor with 25% conditional branches, 80% of which are taken. It is assumed that instructions), and there is no branch prediction. Indicate the correct calculation of CPI:
- a. CPI = 1 + 0.8 \* 0.25 \* 1
  - b. CPI = 1 + 0.8 \* 0.25 \* 2
  - c. CPI = 1 + 0.25 \* 0.8 \* 3
  - d. CPI = 1 + 0.25 \* 0.8 \* 4
13. What is the latency of a branch instruction?
- a. The number of cycles between the IF stage and the MEM stage.
  - b. The number of additional cycles between the IF stage and the cycle of its complete resolution.
  - c. The number of additional cycles between the ID stage and the cycle of its complete resolution.
  - d. The number of cycles between the IF stage and the LD stage.
14. Consider the following source code after applying software speculation:
- ```
LD    R1, 1(R3) ; load A
LD    R1, 0(R2) ; speculative load of B
BEQZ R1, L3 ; branch
DADDI R1, R1, #4 ; else clause
L3: SD    R1, 0(R3) ; speculative store of A
```
- Indicate which of the following statements is correct:
- a. Software speculation is incorrect because the registers used by the speculative load of B have not been renamed.
  - b. Software speculation is correct because the speculative load of B may not succeed.
  - c. Software speculation is incorrect because the "else" clause should be executed before the speculative load of B.
  - d. Software speculation is correct since the branch instruction will execute before the speculative load of B succeeds or fails.

15. Consider the following source code and assume that exception handling is performed using poisoned bits:

```

LD    R1, 1(R3)      ; load A
LD    R1, 0(R2)      ; speculative load of B
BEQZ  R1, L3         ; branch
L3:   DADDI R1, R1, #4 ; else clause
      SD    R1, 0(R3)  ; speculative store of A

```

The state of the registers before executing this code is:

R1	R2	R3	R4	R5	R6	R7	R8	R9
10	0	5	4	5	6	7	8	9

Furthermore, access to memory address 0 is prohibited, so any attempt to read from or write to that position will raise an exception for illegal access. Indicate which of the following statements is correct.

- a. The load instruction for A generates an exception, and the program terminates.
- b. The speculative load of B generates an exception, sets the poisoned bit in register R1 to 1, and the program continues correctly.
- c. The speculative load of B generates an exception, sets the poisoned bit in register R1 to 1, but the program terminates when the branch instruction tries to access that register.
- d. The speculative load of B generates an exception, sets the poisoned bit in register R1 to 1, and the program executes correctly, but the program terminates when the speculative store instruction tries to use register R1.

16. In a processor that uses the Tomasulo technique with speculation, indicate which of the following statements is correct:

- a. In the results' write stage, operations send their results to the reservation stations, and from there, they are sent to the CDB.
- b. In the execution stage, the operation is executed even if the operands are not available, as the CDB will provide them.
- c. In the commit stage, any out-of-order instruction can be completed as long as it is in the reorder buffer.
- d. In the issue stage, if a reservation station is available and there is space in the reorder buffer, the processor continues the execution of an instruction.

17. A difference between a superscalar processor and a VLIW one is...

- a. Conceptually there is no difference. "VLIW" is the name Intel gives to its superscalar processors.
- b. Unlike VLIWs, superscalars can have multiple instructions in the execution stage (EX).
- c. In VLIWs, instruction execution is always in order, while in superscalars it can also be out of order.
- d. None of the above.

18. A 4-way superscalar processor...

- a. Can issue 4 instructions per cycle.
- b. It has a speedup of 1/4 with respect to a scalar processor.
- c. It has a CPI of 4 cycles.
- d. None of the above.

19. In a typical superscalar processor, branch instructions are predicted...

- e. In the Fetch step
- f. In the Decode step
- g. In the Execution step
- h. In the Commit step

20. In a typical superscalar processor, in the case of a misprediction...

- a. An exception occurs that stops program execution.
- b. Program execution continues as normal along the predicted path
- c. The results of the instructions (stored temporarily in the ROB) are copied to the corresponding registers and the execution of the program continues along the alternative path to the one predicted.
- d. None of the above.

21. We have the following optimized scheduling for a 2-way superscalar processor:

- 29. If a thread is running until some event generates a delay, this means that the processor is multi-threading of:
  - a. Fine-grained or interleaved
  - b. Coarse-grained or blocked
  - c. Simultaneous or SMT
  - d. None of the above is true
- 30. Today in the TOP500 most computers are:
  - a. Cluster
  - b. SIMD
  - c. SIMD
  - d. MISD

Cycle	Integer instruction	Fp instruction
1	L.D F1, O(R1)	
2	L.D F8, -8(R1)	MUL.D F4, F2, F0
3	L.D F14, -16(R1)	MUL.D F10, F8, F0
4	L.D F6, O(R2)	MUL.D F16, F14, F0
5	L.D F12, -8(R2)	
6	L.D F18, -16(R2)	
7	DADDI R1, R1, #24	
8	DCONCAT R2, R2, R2, R2	ADD.D F6, F4, F6
9	SQRT R3, R1, end	ADD.D F12, F10, F12
10		ADD.D F18, F16, F18
11	S.D F6, 24(R2)	
12	S.D F12, 16(R2)	
13	BSEQZ R3, Loop	
14	S.D F18, 8(R2)	

Furthermore, we can state that the execution step of the multiplication instruction takes 2 cycles.

b. We can state that the execution step of the multiplication instruction takes 3 cycles.

c. We can state that the execution step of the multiplication instruction takes 4 cycles.

d. We can state that the execution step of the multiplication instruction takes 5 cycles.

22. We must schedule a sequential code consisting of 12 instructions for a VLIW processor with 6 operations per instruction. In the best case (assuming an ideal code and processor) we would have

- a. 2 long instructions.
- b. 3 long instructions.
- c. 4 long instructions.
- d. 6 long instructions.

23. When we improve integration technology:

- a. We quadratically increase the number of transistors/cm<sup>2</sup>
- b. We increase the heat to dissipate quadratically
- c. We increase the power dissipation quadratically
- d. All of them are true

24. Neuvic ARM processors that apply to:

- a. Low power consumption for microcontrollers and IoT
- b. Real-time for critical applications
- c. HPC
- d. Applications for computationally complex systems

25. The Big Little ARM configuration is for:

- a. Having an integrated system for 32-bit and 64-bit applications on the same chip
- b. Have the possibility to address memory in two different formats
- c. Having two types of cores on the same chip
- d. All of them are false

26. The M2 processor has

- a. 8 cores
- b. RISC-V architecture
- c. Tri-Cluster configuration
- d. All of them are false

27. Which of the following is true?

- a. CLOUD is a collection of computers that belong to several owners in different locations and are connected together so that users can share the combined power of the resources.
- b. GRID is a collection of computers usually owned by a single party
- c. Facebook is an example of CLOUD defined as PaSS (Platform as a Service)
- d. None are true

28. A message passing environment that provides tools for communicating processes is:

- a. CUDA

- b. OpenMP

- c. MPI

- d. None of them are true

29. If a thread is running until some event generates a delay, this means that the processor is multi-threading of:

- a. Cluster
- b. SIMD
- c. SIMD
- d. MISD

Name: \_\_\_\_\_

**Important:** Solve each exercise in a different sheet. You can use more sheets if you need, but do not forget to write down your name in all the sheets you deliver.

**Exercise 1 (2 points).** Let's consider the execution of the first 6 iterations of the following loop on a MIPS64 processor:

```
DADD R1', R0, R0      ; (B1)
L1: BEQZ R1, L2      ; (B2)
...                   ;
L2: XORI R1, R1, #1 ; (<-> NOT R1)
J   L1

```

- Indicate the sequence of branch instructions (B1/B2) executed and their actual outcomes (T/N/T). Use the first and second columns in the table below.
- Detail the behavior of a one-level branch predictor using 2 bits per branch instruction. Consider the initial state indicated in the table. Use columns 3, 4, and 5. Indicate at the end the hit rate for this predictor.
- Detail the behavior of a two-level (2, 2) correlating branch predictor, assuming the same input vector, and considering the initial state indicated in the table. Use columns 6, 7, and 8. Indicate at the end the hit rate for this predictor.

c) Detail the execution of the first 6 iterations of the following loop on a MIPS64 processor:

a) Actual behavior		b) One-level predictor behavior		c) Two-level predictor behavior	
Branch instruction executed (B1/B2)	Actual outcome (T/N/T)	Prediction bits for B1 and B2	Predicted outcome (T/N/T)	Prediction bits (BHR and BTHs)	Predicted outcome (T/N/T)
(Initial state)	-	00	NT	N	"
B1	T	01 02	NT	00 00	"
B2	T	02 01	NT	11 " 01	"
B1	NT	00 04	NT	10 "	00
B2	NT	00 05	NT	01 00	"
B1	T	04 00	NT	11 " 00	"
B2	T	05 01	NT	10 "	00
B1	NT	01 00	NT	11 " 00	"
B2	NT	00 00	NT	10 "	00
B1	T	00 01	NT	11 " 00	"
B2	T	01 02	NT	01 00	"
B1	NT	01 00	NT	11 " 00	"
B2	NT	02 00	NT	10 "	00

Hit rate: **50%**

Hit rate: **66'66%**

$$\begin{aligned} &= 0.5 \cdot 100 \\ &= 50\% \end{aligned}$$

$$\begin{aligned} &= 0.6666 \cdot 100 \\ &= 66'66\% \end{aligned}$$

Name: \_\_\_\_\_

**Important:** Solve each exercise in a different sheet. You can use more sheets if you need, but do not forget to write down your name in all the sheets you deliver.

→ Pa branches

**Exercise 2 (3 points).** The following MIPS64 code is going to be executed on a processor without speculation but with dynamic scheduling using the Tomasulo technique.

```
L.D    F2, 0(R1)      /* I1 */ 1F D EX MEM WB
L.D    F4, 8(R1)      /* I2 */ 1F D EX MEM WB
L.D    F6, 16(R1)     /* I3 */ 1F D EX MEM WB
ADD  D F8, F2, F4   /* I4 */ 1F D EX MEM WB
MUL  D F10, F2, F6  /* I5 */ 1F D EX MEM WB
DIV  D F12, F8      /* I6 */ 1F D EX MEM WB
S.D    F13, 32(R1)    /* I7 */ 1F D EX MEM WB
```

The mentioned processor has the following characteristics:

- 1 cycle for loads and stores.
- 4 cycles for floating-point multiplication (FP).
- 3 cycles for floating-point addition (FP).
- 6 cycles for floating-point division (FP).
- Reservation stations per FP unit: 2 entries.
- Load buffer: 2 entries.
- Store buffer: 1 entry.
- The execution stage of an instruction waiting for a result cannot start in the same cycle in which the result is written to the CDB (Common Data Bus).

Please answer the following questions for the given code and the provided processor configuration.

- a) Identify the dependencies between instructions and indicate their type in the following table:

Instruction 1	Instruction 2	Type of dependency
		RAW

6/12 =	
= 0.5 · 100	Hit rate: <b>50%</b>

8/P.

$$\begin{aligned} &= 0.6666 \cdot 100 \\ &= 66'66\% \end{aligned}$$

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING	
COMPUTING SYSTEMS DEPARTMENT	
Final written exam (June 2023)	
UCLM	E <i>Scuela</i> Ingeniería Informática
Name:	

- b) Indicate the stall cycles between instructions, if any, in the gaps in the previous code that appear in the following table:

L.D	F2, 0 (R1)	0
L.D	F4, 8 (R1)	0
L.D	F6, 16 (R1)	0
ADD.D	F8, F2, F4	2
MUL.D	F10, F2, F6	5
DIV.D	F0, F12, F8	5
S.D	F0, 32 (R1)	8

- c) Fill in the following table using the Tomasulo technique:

cycle	Mem 1		Mem 2		FP 1	FP 2	integer/branch
	Issue	Execution	Commit	Issue	Execution		
1	L.D F2, 0 (R1)	L.D F8, -8 (R1)	NOP	NOP	NOP	NOP	NOP
2	L.D F14, -16 (R1)	L.D F20, -24 (R1)	NOP	NOP	NOP	NOP	NOP
3	L.D F6, 0 (R2)	L.D F12, -8 (R2)	MUL.D F4, F2, F0	MUL.D F10, F8, F0	MUL.D F22, F20, F0	NOP	NOP
4	L.D F18, -16 (R2)	L.D F24, -24 (R2)	NOP	NOP	NOP	NOP	NOP
5	NOP	NOP	NOP	NOP	NOP	NOP	NOP
6	NOP	NOP	NOP	NOP	NOP	NOP	DADDUI R1, R1, #32
7	NOP	NOP	ADD.D F6, F4, F6	ADD.D F12, F10, F12	NOP	NOP	NOP
8	NOP	NOP	ADD.D F18, F16, F16	ADD.D F24, F22, F24	NOP	NOP	NOP
9	NOP	NOP	NOP	NOP	NOP	NOP	SLT1 R3, R1, \$1
10	NOP	NOP	NOP	NOP	NOP	NOP	DADDUI R2, R2, #32
11	S.D F6, 0 (R2)	S.D F12, -8 (R2)	NOP	NOP	NOP	NOP	NOP
12	S.D F18, 16 (R2)	S.D F24, 8 (R2)	NOP	NOP	NOP	NOP	BEQZ R3, Loop 1

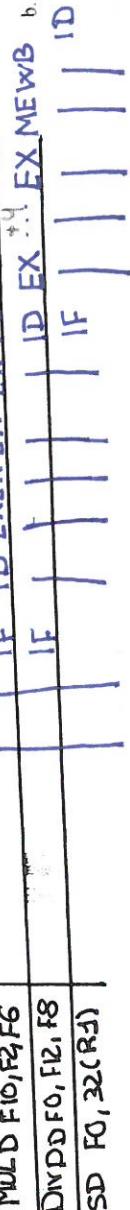
Assume that the code is as compact as possible, and the processor employs forwarding and delayed branch.

- a. Indicate the cycles used by the functional units (FU) of the processor and complete the following table:

Instruction	Issue	Execution	Commit	Functional unit	Number of cycles
L.D F2, 0 (R1)	1	2	3	Integer FU	1
L.D F4, 8 (R1)	2	3	4	Adder FP FU	2
L.D F6, 16 (R1)	3	4	5	Multiplexer FP FU	3
ADD.D F8, F2, F4	4	5-7	8		4
MUL.D F10, F2, F6	5	6-9	10		4
DIV.D F0, F12, F8	6	9-14	15		6
S.D F0, 32 (R1)	7	16	17		2

- b) Show the initial code for the MIPS64 processor segmented such that at each turn of the loop a result is calculated. Write the code in this table (use as many rows as you need).

Inst: 1	LD F2, 0(R1)
Inst: 2	LD F6, 0(R2)
Inst: 3	MULD F4, F2, F0
Inst: 4	ADD.D FG, F4, F6
Inst: 5	DADDI R1, R1 + 8
Inst: 6	SLTI R3, R2, F1
Inst: 7	DADDUI R2, R2, #32



Inst. 8	SD F6,O(R2)
Inst. 9	BEGZR3,loop
Inst. 10	
Inst. 11	
Inst. 12	
Inst. 13	
Inst. 14	
Inst. 15	

- c. What is the speedup obtained by the VLIW processor with respect to the code in section b) of the segmented MIPS4 processor? Explain the calculations made and consider that the code is scheduled to obtain the best performance (i.e. the lowest number of stalls).

12 codes / iteration VLIW



Name: \_\_\_\_\_ Test Year 2020/21

Aula 1.5. Mesa número: \_\_\_\_\_

Indicate the correct answer to each test question in the next table:

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23	24	25	26	27	28	29	30

**Important:** Test counts **3 points**. Each correct question counts 0.1 points. Wrong answers subtract 0.13 points. Not answered questions are not considered for the mark.

1. Two examples of processors with the same organization but different architecture are:

- a) Itanium and Pentium
- b) AMD Opteron and Pentium
- c) ARM and Alpha
- d) None of the above

2. MIPS allow us to compare the performance of:

- a) Computers with the same instruction set
- b) Computers with different instruction set
- c) Compilers
- d) All the above

3. We want to compare the performance of two different systems: S1 and S2. S1 costs 1000 € and S2 costs 1200 €.

The following times have been obtained in these systems:

Program	S1 time	S2 time
1	30 s	26 s
2	8 s	7 s

Which machine is more cost effective and how much?

- a) S1 is 1.042 more cost effective than S2
- b) S1 is 1.14 more cost effective than S2
- c) S2 is 1.14 more cost effective than S1
- d) None of the above

4. We have a computer with a clock frequency of 1000MHz, requiring the following number of clock cycles to execute each type of instruction:

- Floating point multiplication: 6 cycles
- Floating point addition: 4 cycles
- Floating point division: 20 cycles
- Integer instructions: 2 cycles

Consider a program (P), with the next instruction mix:

- Floating point multiplication: 10%
- Floating point addition: 15%
- Floating point division: 5%
- Integer instructions: 70%

The native MIPS value of this computer is:

- a) 277.78
  - b) 321.3
  - c) 211.1
  - d) None of the above
5. We are looking for a benchmark to measure the performance of the new floating-point unit of a computer. This unit runs 4 times faster than the old one. We want to get a speedup of 2 for the benchmark we are considering. If the execution time of this benchmark in the old floating-point hardware is 100 seconds, how much of that initial execution time should the floating-point instructions count to get this speedup?
- a) 80 s
  - b) 66.6 s
  - c) 50 s
  - d) None of the above
6. Given two instructions *i* and *j* (*i* precedes *j*), there is a data dependence between them if:
- a) Both instructions update the same register or memory position.
  - b) Instruction *i* produces a result that may be used by instruction *j*.
  - c) Instruction *j* produces a result that may be used by instruction *i*.
  - d) Instruction *j* writes a register or memory position that instruction *i* reads.
7. In this loop:
- ```

Loop: L.D    F0, 0(R1) ;I1
      MUL.D F2, F0, F1 ;I2
      SUB.D F3, F2, F0 ;I3
      S.D   F3, 0(R1) ;I4
      DADD.U R1, R1, #-8 ;I5
      BNEZ  R1, Loop
  
```
- Between I1 and I5 there is...
- a) A data dependence between iterations
  - b) A data dependence in the same iteration
  - c) An antidependence between iterations
  - d) An output dependence in the same iteration
8. Consider the next high-level code (A and B are two previously defined arrays):
- ```

for(i=0; i<100; i++) {
    A[i+1] = A[i] + B[i];
    B[i]   = A[i] * 2;
}
  
```
- a) There is an antidependence, between S1 and S2, through B (between iterations)
- b) There is a data dependence, between S1 and S2, through A (in the same iteration)
- c) There is an antidependence, between S1 and S1, through A (between iterations)
- d) None of the above
9. Output dependences:
- a) Are removed by software pipelining
  - b) Are removed by register renaming
  - c) Are removed by trace scheduling
  - d) None of the above
10. We unroll 3 times this loop (assuming that the number of iterations is divisible by 3):
- ```

Loop: L.D    F0, 0(R1) ;I1
      ADD.D F2, F0, F1 ;I2
      SUB.D F3, F2, F0 ;I3
      S.D   F3, 0(R1) ;I4
      DADD.U R1, R1, #-8 ;I5
      BNEZ  R1, Loop
  
```
- a) The resulting loop will be composed of 6 instructions.
- b) The resulting loop will be composed of 9 instructions.
- c) The resulting loop will be composed of 14 instructions.
- d) The resulting loop will be composed of 18 instructions.
11. After unrolling 3 times the loop of the previous question, the immediate value of the DADD.U instruction will be:
- a) #8
  - b) #-8
  - c) #24
  - d) #-24

12. In Scoreboard, the order in which the issue step is performed is given by:

- a) The availability of functional units.
- b) The order of the instructions in the program.
- c) The antidependencies in the program.
- d) None of the above.

13. To implement the BHR of a (2,2) predictor we use:

- a) A binary adder
- b) A saturating counter
- c) A shift register
- d) None of the above

14. A (2,4) correlating predictor uses exactly:

- a) 2 BHTs with 4 entries
- b) 2 BHTs and 4 BHTs
- c) 4 BHTs and 2 BHTs
- d) None of the above

15. Our processor uses a 2-bit predictor, initialized to 11, for each branch. What is the number of mispredictions for this code?

```
for (i=0; i<10; i++) { branch B1  
    for (j=0; j<10; j++) { ; branch B2  
        . . .  
    }  
}
```

- a) 9
- b) 10
- c) 11
- d) None of the above

16. Compute the CPI of a code in which 30% of the instructions are conditional branches. The code is executed in a 5-stage pipelined processor without BTB, in which the branch target address is known at the end of the MEM stage and the branch outcome (taken or not taken) is known at the end of the EX stage. Assume that 75% of the conditional branch instructions are taken, and a CPI=1 in all cases, except for the taken branches.

- a) CPI=1.675
- b) CPI=1.6
- c) CPI=1.45
- d) CPI=1.225

17. A speculative instruction:

- a) Is an instruction that removes stalls due to control hazards.
- b) Is an instruction whose execution depends on a previously established condition.
- c) Is an instruction causing a terminating or resumable exception.
- d) None of the above.

18. What will be the behavior of a MIPS64 processor that executes the speculative instruction "ADD.D\* F6, F7, F8", assuming the "poison bits" technique and knowing that the poison bits of F7 and F8 are initially 0 and 1, respectively?

- a) The poison bit of F8 will be reset (poison bit=0)
- b) The poison bit of F7 will be set (poison bit=1)
- c) The poison bit of F6 will be reset (poison bit=0)
- d) None of the above

19. A superscalar processor of degree  $n$ :

- a) Issues a variable number of instructions per cycle, which is always less or equal to  $n$ .
- b) Limits to  $n/2$  the number of memory accesses that can be simultaneously issued.
- c) Improves the performance, by reducing the number of instructions executed.
- d) Has a reduced hardware cost.

20. In the *Issue* step of a un superscalar processor:

- a) The instruction is decoded and sent to the instructions window.
- b) The processor determines which instructions in the window can be sent to execution.
- c) The execution of the instruction is carried out in the corresponding functional unit.

21. Resources and operands are checked.

- a) Dynamic renaming.
- b) Recovery in case of misprediction.
- c) Instruction reading bandwidth.
- d) Resources and operands checking.

22. Among the advantages of VLIW processors, we could highlight:

- a) They perform a dynamic scheduling.
- b) Hardware detects dependencies, since compiler is very simple.
- c) Hardware is very simple, since complexity resides in the compiler.
- d) They do not have to unroll the code.

23. The power dissipated by a processor is proportional to:

- a) Frequency
- b) Voltage squared
- c) Transistor count
- d) All the above

24. A difference between the P6 and P7 architectures is:

- a) P6 uses more decoders
- b) P7 does not use barrel shifter
- c) L1 data cache in P6 is bigger
- d) All the above

25. One of the advantages of a multicore over a single core is:

- a) Less PCB than SMP, which implies less energy
- b) A better use of chip area
- c) The share the same bus and memory bandwidth
- d) All of them are disadvantages

26. Indicate the correct statement:

- a) Quantum computers will replace general purpose processors in the medium term.
- b) Heterogeneous computing is present in many of today's mobile phones.
- c) Heterogeneous computing is only applicable to large supercomputers.
- d) None of the above.

27. In a multi-threading processor, several instructions from different threads are issued to be executed at the same time when the type of multi-threading is:

- a) Fine-grained multi-threading
- b) Coarse-grained multi-threading
- c) Simultaneous multi-threading
- d) Chip multiprocessing

28. GRID services include:

- a) Resource discovering and monitoring
- b) Resource allocation and management
- c) Security infrastructure and file transfer
- d) All the above

29. Indicate the correct statement:

- a) A SIMD architecture is less energy efficient than a MIMD
- b) MMX extensions of Intel processors are MISD extensions
- c) In GPUs there is not penalty due to changes of context
- d) None of the above

30. Skype or Gmail are examples of:

- a) Cloud SaaS
- b) Cloud PaaS
- c) Cloud IaaS
- d) Cloud JaaS

Name: \_\_\_\_\_ **Important:** Solve this exercise on this sheet, and do not forget to write down your name in all the sheets you deliver.

**Exercise 2 (2 points).** We have a MIPS64 processor equipped with hardware speculation according to the “Tomasulo” method. The processor has the following characteristics:

- There is an 8-entry ROB (initially EMPTY), 1 load buffer, 1 reservation station (RE) for FP addition, 2 reservation stations for FP multiplication, and 2 reservation stations for FP division.

| Functional Unit | Amount | Cycles | Pipelined |
|-----------------|--------|--------|-----------|
| FP Adder        | 1      | 2      | Yes       |
| FP Multiplier   | 1      | 4      | Yes       |
| FP Divider      | 1      | 8      | No        |
| LOAD/STORE      | 1      | 2      | Yes       |
| INT ALU         | 1      | 1      | No        |

- In the same clock cycle in which a result is written in the Common Data Bus (CDB) the execution of any instruction that was waiting for that result is initiated.

- A resource that is released freed in a clock cycle cannot be used until the next cycle.

- a) Complete the following table, indicating the cycles in which each stage is carried out.

| Instruction            | Issue | Execute | Write result | Commit |
|------------------------|-------|---------|--------------|--------|
| 1.- MUL.D F2, F1, F0   | 1     | 2 - 5   | 6 (L1,b)     | 7      |
| 2.- DIV.V D F0, F2, F4 | 2     | 6 - 13  | 15           | 15     |
| 3.- ADD.D F2, F4, F5   | 2     | 7 - 5   | 9            | 16     |
| 4.- S.D E2, 0(R1)      | 4     | 6 - 3   | 8            | 17     |
| 5.- MUL.D F4, F2, F6   | 7     | 8 - 11  | 12           | 18     |
| 6.- DIV.D F6, F1, F2   | 15    | 16 - 23 | 24           | 23     |
| 7.- L.D F6, 0 (R2)     | 25    | 26 - 23 | 28           | 29     |
| 8.- ADD.D F2, F0, F6   | 28    | 29 - 30 | 31           | 32     |
| 9.- MUL.D F3, F0, F2   | 31    | 32 - 35 | 36           | 37     |
| 10.- ADD.D F3, F0, F6  | 36    | 37 - 38 | 39           | 40     |

start-up: LW RS, 0(R1)

loop:

SW RS, 0(R1)

DADD R3, R3, RS

SLT RS, R4, RS

SLT RS, R3, RS

DADD R3, R3, RS

SW RS, 8(R2)

DADD R3, R3, RS

SW RS, 8(R2)

0, 4, 8, 12

finish-up:

SW RS, 0(R1)

DADD R3, R3, RS

SLT RS, R4, RS

SLT RS, R3, RS

DADD R3, R3, RS

SW RS, 4(R2)

DADD R3, R3, RS

SW RS, 4(R2)

0, 4, 8, 12

| Stall due to             | FU                 | Where? |
|--------------------------|--------------------|--------|
| Struct hazard due to FU  | Occupied by DIV.D2 |        |
| Struct hazard due to ROB |                    |        |
| Data hazard due to FU    |                    |        |
| Struct hazard due to RS  |                    |        |

| Stall due to             | FU                 | Where? |
|--------------------------|--------------------|--------|
| Struct hazard due to FU  | Occupied by DIV.D2 |        |
| Struct hazard due to ROB |                    |        |
| Data hazard due to FU    |                    |        |
| Struct hazard due to RS  |                    |        |

| Stall due to             | FU                 | Where? |
|--------------------------|--------------------|--------|
| Struct hazard due to FU  | Occupied by DIV.D2 |        |
| Struct hazard due to ROB |                    |        |
| Data hazard due to FU    |                    |        |
| Struct hazard due to RS  |                    |        |

| Stall due to             | FU                 | Where? |
|--------------------------|--------------------|--------|
| Struct hazard due to FU  | Occupied by DIV.D2 |        |
| Struct hazard due to ROB |                    |        |
| Data hazard due to FU    |                    |        |
| Struct hazard due to RS  |                    |        |

| SCHOOL OF COMPUTER SCIENCE AND ENGINEERING     |
|------------------------------------------------|
| COMPUTING SYSTEMS DEPARTMENT                   |
| COMPUTER SCIENCE DEGREE, COMPUTER ARCHITECTURE |

Final exam (January 2021)      Exercises      Year 2020/21

Name: \_\_\_\_\_  
**Important:** Solve this exercise on this sheet, and do not forget to write down your name in all the sheets you deliver.

**Exercise 1 (2.5 points).** We have installed a CO<sub>2</sub> monitor in our classroom. Every 5 minutes, it gets a new sample of the CO<sub>2</sub> level. Each sample is a 32-bit positive value. All these samples are stored in a set of consecutive memory positions (starting from the address in R1). Assuming that R2 contains the number of samples taken during a whole day, the following MIPS64 code counts (in R3) the number of times the CO<sub>2</sub> level has exceeded a certain threshold (in R4), resets all the samples below that threshold and writes 1 in the rest of them. Assume that R1, R2, R3 and R4 have been previously initialized.

```

    LW      R5, 0(R1)        #I1 R5 = next CO2 sample
    SLT    R5, R4, R5       #I2 checks if it exceeds the threshold (R4)
    DADD   R3, R3, R5       #I3 updates the counter (R3)
    SW      R5, 0(R1)        #I4 overwrites the sample
    DADDUI R1, R1, #4       #I5 R1 = address of the next CO2 sample
    DADDUI R2, R2, -1       #I6 updates R2
    BNE    R2, R0, read      #I7 If R2>>0 read the next CO2 sample
    
```

Assume that we have a pipelined processor with following characteristics:

- Operand forwarding is supported
- A value can be written and read from a register in the same clock cycle
- All the dependences are detected and solved in the ID stage
- Branches are completely solved in the ID stage
- Delayed branch is NOT supported

a) Indicate the stalls in the code between each pair of instructions:

```

    LW      R5, 0(R1)        #I1 R5 = next CO2 sample
    stalls → NO
    SLT    R5, R4, R5       #I2 checks if it exceeds the threshold (R4)
    stalls → LW, data → 1Stall
    DADD   R3, R3, R5       #I3 updates the counter (R3)
    stalls → NO, Forwarding → EX→EX
    SW      R5, 0(R1)        #I4 overwrites the sample
    stalls → NO MODIFICA
    DADDUI R1, R1, #4       #I5 R1 = address of the next CO2 sample
    stalls → NO MODIFICA
    DADDUI R2, R2, -1       #I6 updates R2
    stalls → NO MODIFICA
    BNE    R2, R0, read      #I7 If R2>>0 read the next CO2 sample
    stalls → 1. Decode on ID
    Next instruction after the BNE
    
```

b) Apply software pipelining and show the modified loop body.

read:

*(Handwritten notes and diagrams are present in the margins, including a large blue circle around 'NO FORWARDING' and arrows pointing to various stages of the pipeline.)*

c) Show the start-up and finish-up blocks (do not schedule them):

Start-up:

Finish-up:

d) Show the complete code (do not schedule it, but remember that the new code must behave exactly like the initial one):

Name: \_\_\_\_\_  
**Important:** Solve this exercise on this sheet, and do not forget to write down your name in all the sheets you deliver.

**Exercise 3 (2.5 points).** Given the following code fragment, where R2 and R3 contain the initial and final memory addresses where the data are stored:

```
Bucle: L.D    F10, 0(R2)
      ADD.D F8, F10, F12
      S.D    F8, 0(R2)
      DADDUI R2, R2, -8
      BNE   R2, R3, Bucle
```

We want to compare the number of instructions issued per cycle by processors A and B, with the following common characteristics:

- We assume forwarding and delayed branch.
- All the units are fully pipelined.
- The next latencies:

| Instruction producing the result | Instruction using the result | Latency in cycles |
|----------------------------------|------------------------------|-------------------|
| ALU FP                           | ALU FP                       | 3                 |
| ALU FP                           | SD                           | 1                 |
| L.D                              | ALU FP                       | 1                 |
| L.D                              | SD                           | 0                 |

Each processor presents the following specific characteristics:

- **Processor A.** It is a MIPS64 5-stage superscalar processor, capable of issuing 4 instructions per cycle: two integer instructions (including memory accesses and branches) and two floating-point instructions.
- **Processor B.** It is a VLIW processor, capable of issuing in each cycle four operations per instruction. It does not impose restrictions on the type of operations that can be issued.

a) For processor A, unroll 8 times the code and schedule it as compact as possible, by indicating the instructions issued in each cycle. Use as rows as you need in the next table:

| Cycle | MEM            |              | FP           |                                  |
|-------|----------------|--------------|--------------|----------------------------------|
|       | Instr. 1       | Instr. 2     | Instr. 3     | Instr. 4                         |
| 1     | L.D F10, 0     | L.D F16, -8  |              | L.D F16, -8 (R2)                 |
| 2     | L.D F20, -16   | L.D F24, -24 | ADD F8, F10  | ADD F14, F16, F18, F20, F24, F12 |
| 3     | L.D F28, -32   | L.D F32, -40 | ADD F18, F20 | ADD F22, F24                     |
| 4     | L.D F4, -48    | L.D F7, -56  | ADD F26, F28 | ADD F30, F31                     |
| 5     | SD F8, 0       | SD F14, -8   | ADD F2, F4   | ADD F6, F7                       |
| 6     | SD F18, -16    | SD F22, -24  |              |                                  |
| 7     | SD F26, -32    | SD F30, -40  |              |                                  |
| 8     | SD F2, -48     | SD F6, -56   |              |                                  |
| 9     | DADDUI R2,R2   |              |              |                                  |
| 10    | BNE R2, R3, R2 | R3,R2        |              |                                  |

→ Par la data dependency, espera 2 más

- b) For processor B, schedule the unrolled code, assuming that static scheduling is as aggressive as possible and as efficient and possible to optimize the resulting code. You can use register renaming and out-of-order scheduling. Use as rows as you need in the next table:

| Cycle | Operation 1 | Operation 2 | Operation 3 | Operation 4 |
|-------|-------------|-------------|-------------|-------------|
| 1     |             |             |             |             |
| 2     |             |             |             |             |
| 3     |             |             |             |             |
| 4     |             |             |             |             |
| 5     |             |             |             |             |
| 6     |             |             |             |             |
| 7     |             |             |             |             |
| 8     |             |             |             |             |
| 9     |             |             |             |             |
| 10    |             |             |             |             |

- c) Compare the number of operations per cycle for each processor. Which processor gets the best efficiency in operations per cycle?
- a)  $\frac{26 \text{ instruc}}{10 \text{ ciclos}} = 2.6 \text{ inst/ciclo}$       b)  $\frac{26 \text{ instruc}}{7 \text{ ciclos}} = 3.70 \text{ inst/ciclo}$
- ↑ ↑ mejor ↑↑



loop: LB R2, 0(R1)  
 DADD R3, R2, R0  
 BNE R2, R4, not\_a  
 DADDUI RS, RS, 1  
 DADDUI R6, R6, 1 ← not\_a:  
 DADDUI R1, R1, 1  
 BNE2 R3, loop

Carga "exam18"

(m,n)  
(2,1)

$2^m$  branches of n bits

| Branch executed (initial state) | Actual outcome | Predicted outcome | Hit/miss | BHR | BHTs |    |    |    |
|---------------------------------|----------------|-------------------|----------|-----|------|----|----|----|
| B1 (e)                          | T              | NT                | miss     | 01  | 00   | 00 | 00 | 00 |
| B2 (e)                          | T              | NT                | miss     | 11  | 10   | 01 | 00 | 00 |
| B1 (x)                          | T              | NT                | miss     | 11  | 10   | 01 | 00 | 10 |
| B2 (x)                          | T              | NT                | miss     | 11  | 10   | 01 | 00 | 11 |
| B1 (a)                          | NT             | T                 | miss     | 10  | 1001 | 00 | 01 |    |
| B2 (a)                          | T              | NT                | miss     | 01  | 1001 | 01 | 01 |    |
| B1 (m)                          | T              | NT                | miss     | 11  | 1011 | 01 | 01 |    |
| B2 (m)                          | T              | T                 | hit      | 11  | 1011 | 01 | 01 |    |
| B1 (1)                          | T              | NT                | miss     | 11  | 1011 | 01 | 11 |    |
| B2 (1)                          | T              | T                 | hit      | 11  | 1011 | 01 | 11 |    |
| B1 (8)                          | T              | T                 | hit      | 11  | 1011 | 01 | 11 |    |
| B2 (8)                          | T              | T                 | hit      | 11  | 1011 | 01 | 11 |    |
| B1 (..)                         | T              | T                 | hit      | 11  | 1011 | 01 | 11 |    |
| B2 ( )                          | NT             | T                 | miss     | 10  | 1011 | 01 | 10 |    |

hitrate:  $5/14 \cdot 100 = 35.75\%$

Most significant bit = el que + posición  
2-bit saturating  $\rightarrow 00, 01, 10, 11$

loop: LW R3, 0(R1)  
 BEQZ R3, next B2  
 DMUL R3, R3, RS  
 SW R3, 0(R1)

Carga: 0, 1, 1, 2, 3, 5

R2  $\rightarrow$  Counter of number

next: DADDI R1, R1, 4  
DADDUI R2, R2, -1

BNE2 R2, loop B2  
Branch executed T/NT Predictor used(P1/P2)

| Branch executed | T/NT | Predictor used(P1/P2) | Predicted T/NT | Hit/miss | Counters |    | Predictor 1 |    | Predictor 2 |     |
|-----------------|------|-----------------------|----------------|----------|----------|----|-------------|----|-------------|-----|
|                 |      |                       |                |          | B1       | B2 | B1          | B2 | B1          | B2  |
| B1 <sub>0</sub> | T    | P2                    | NT             | miss     | 00       | 00 | 0           | 0  | 000         | 000 |
| B2              | T    | P2                    | NT             | miss     | 00       | 00 | 1           | 1  | 001         | 001 |
| B1 <sub>1</sub> | NT   | P2                    | NT             | hit      | 00       | 00 | 0           | 1  | 000         | 001 |
| B2              | T    | P2                    | NT             | miss     | 00       | 01 | 0           | 1  | 000         | 010 |
| B1 <sub>2</sub> | NT   | P2                    | NT             | hit      | 00       | 01 | 0           | 1  | 000         | 010 |
| B2              | T    | P2                    | NT             | miss     | 00       | 10 | 0           | 1  | 000         | 011 |
| B1 <sub>3</sub> | NT   | P2                    | NT             | hit      | 00       | 10 | 0           | 1  | 000         | 011 |
| B2              | T    | P1                    | T              | hit      | 00       | 11 | 0           | 1  | 000         | 100 |
| B1 <sub>4</sub> | NT   | P2                    | NT             | hit      | 00       | 11 | 0           | 1  | 000         | 101 |
| B2              | NT   | P1                    | T              | miss     | 00       | 11 | 0           | 0  | 001         | 100 |

| Instruc      | Issue | Execute | Write | Commit |
|--------------|-------|---------|-------|--------|
| LD F0, 0(R1) | 1     | 2 - 3   | 4     | 5      |
| LD F2, 0(R1) | 2     | 3 - 4   | 5     | 6      |
| LD F4, 0(R3) | 5     | 6 - 7   | 8     | 9      |

Pq este salto? pq hay 2 RS (Reservation stations) por FU. Así se veía al llegar a la 3<sup>a</sup> instruc  $\Rightarrow$  RS1 Busy | RS2 Busy. Ambas ocupadas.

MULD F6, F0, F2 6 7 - 12 13 14

Hay data hazards pero no afecta pq la instruc 3 hizo stalls por su problema. Si llega a haber 3 RS de Mem (LD o SD), Mulp debería haber esperado igualmente al ciclo 6 para ejecutar que es cuando F2 tiene el resultado. Esperaría al 6 ya que tiene que esperar al siguiente clock cycle del que escribe.

ADD F8, F0, F4 7 9 - 10 11 15  $\rightarrow$  por el ROB debe acabar en orden

¿Y el ciclo 8?  $\rightarrow$  No puede ejecutar pq es ahí cuando LD F4, 0(R3) escribe y

DIV.D F10, F8, F6 8 14 - 23 24 25

¿Este salto?  $\rightarrow$  por data dependency con MULD F6, F0, F2 en F6

SD F10, 0(R4) 9 25 - 26 27?? 28

¿Salto? Data dependency  
con DIV.D F10, F8, F6 en F10

$\hookrightarrow$  pq las "????" los stores no escriben en memoria, ACCEDEN, pero en clase lo tomamos como escritura.

| FU    | FU | Cycles | Instuc          | Issue    | Read     | Exec    | Write      |
|-------|----|--------|-----------------|----------|----------|---------|------------|
| ADD   | 1  | 4      | LD F0, 0(R1)    | 1        | 2        | 3       | 4          |
| MULT  | 1  | 6      | LD F2, 4(R1)    | 5 (Edu)  | 6        | 7       | 8          |
| DIV   | 1  | 8      | LD F4, 8(R1)    | 9 (Edu)  | 10       | 11      | 12         |
| LD/SD | 1  | 1      | MULD F6, F0, F2 | 10       | 11       | 12 - 17 | 18         |
| Int   | 1  | 1      | MULD F4, F8, F6 | 19 (Edu) | 20       | 21 - 26 | 27         |
|       |    |        | ADDD F8, F0, F8 | 20       | 21       | 22 - 25 | 28 (orden) |
|       |    |        | SD F4, 0(R4)    | 21       | 28 (RAW) | 29 ??   | 30         |

$\rightarrow$  Regla de bloqueo:  
cycle + 1

$\rightarrow$  Para no sobreescribir

$\hookrightarrow$  MAL, por esto pipelada pondrá estas ambas a la vez, siempre y cuando NO sea en la misma etapa.

Tomasito

## Reservation Station

| Name | $Q_j$ | $Q_k$ | $V_j$ | $V_k$ | Address | Data | Operation |
|------|-------|-------|-------|-------|---------|------|-----------|
|------|-------|-------|-------|-------|---------|------|-----------|

DIVD F4, F2, FO  
Add F6, F4, F3  
J K

(Supongamos fin  
in-order)

Sí, órdenes en Register Result Status  
pueden ser fijas de orden

3 LD stat.

3 ADD stat.

2 MUL stat.

LD → 2 cycle

ADD → 2 "

MUL → 10 "

DIV → 40 "

FU

FO F2 F4 ... F3c

| Instruc         | ISSUE | Execute | Write | Commit |
|-----------------|-------|---------|-------|--------|
| LD F6, 34(R2)   | 1     | 2 - 3   | 4     | 5      |
| LD F2, 45(R3)   | 2     | 3 - 4   | 5     | 6      |
| MULT F0, F2, F4 | 3     | 6 - 15  | 16    | 17     |
| SUB F8, F6, F2  | 4     | 6 - 7   | 8     | 18     |
| DIV F10, F0, F6 | 5     | 17 - 56 | 57    | 58     |
| ADD F6, F8, F2  | 6     | 19 - 20 | 21    | 59     |

RS

| Time | Name  | Busy | OP   | $V_j$  | $V_k$ | $Q_j$ | $Q_k$                                   |
|------|-------|------|------|--------|-------|-------|-----------------------------------------|
|      | ADD1  | Yes  | SUB  | 34(R2) |       | Load2 | { Paro en ciclo 4}                      |
|      | ADD2  | Yes  | ADD  |        | LD2   | ADD1  |                                         |
|      | ADD3  |      |      |        |       |       |                                         |
|      | MULT1 | Yes  | MULT |        | F4    | Load2 | { Paro en ciclo 3, Qj: esperando de LD} |
|      | MULT2 | Yes  | Div  |        | F6    | Mult1 | { Paro en ciclo 5}                      |

### Mini-exercises

2) Benchmark executed on 40MHz processor

$\rightarrow f = \text{clock rate}$

| Instruc type | inst Count |
|--------------|------------|
| Arith        | 45000      |
| Branch       | 32000      |
| LD/SD        | 15000      |
| FP           | 8000       |

| clock cycle |
|-------------|
| 1           |
| 2           |
| 2           |
| 2           |

$$\text{Clock speed} = 40 \text{ MHz} = 40 \cdot 10^6 \text{ Hz}$$

b) Exec. time?

$$\frac{\text{Total cycles}}{\text{freq}} = \frac{155000}{40 \cdot 10^6} = 0'003875 \text{ s}$$

a) Avg CPI?

$$\frac{\text{Cycles to execute}}{\text{Instruc. Count}} =$$

$$= \frac{45000 \cdot 1 + 32000 \cdot 2 + 15000 \cdot 2 + 8000 \cdot 2}{45000 + 32000 + 15000 + 8000} = 1'55$$

~ 155 000  
~ 100 000

c) MIPS rate?

$$\frac{\text{Instruction Count}}{\text{Exec time}} = \frac{100000}{0'003875} = 251'8 \text{ MIPS}$$

c) Throughput?

$$\frac{\text{MIPS} \cdot 10^6}{\text{Instruc Count}} = \frac{251'8 \cdot 10^6}{100000} = 258$$

$$\frac{f}{\text{Instruc Count} \cdot \text{CPI}} = \frac{40 \cdot 10^6}{100000 \cdot 1'55} = 258'06 \dots$$

3) 100000 instruc on 500 MHz processor

| Instruc type | CPI | instruc % |
|--------------|-----|-----------|
| Arith        | 1   | 60%       |
| FP           | 2   | 20%       |
| LD/SD        | 4   | 10%       |
| MEM          | 6   | 10%       |

a) AVG CPI =  $1 \cdot 0'6 + 2 \cdot 0'2 + 4 \cdot 0'1 + 6 \cdot 0'1 = 2 \text{ cycles/instr}$

b) MIPS =  $\frac{f}{\text{CPI}} = \frac{500 \cdot \text{MHz}}{2} = 250$

## VLIW and Superscalers Exercises

Int: 1 FPDIV = 16  
FP mult: 5

1) Superscalar issues 4 instruc per cycle. 2 integer, 2 FP. FP ADD  $\Rightarrow$  3 cycles per exec

loop:  
 LD F0,0(R1)  
 ADD.D F4,F0,F2  
 SD F4,0(R1)  
 DADDI R1,R1,-8  
 BNEZ R1,loop

a) Unroll 8 times, schedule as compact as possible

LD F0,0(R1)  
 ADD F0,F0,F2  
 SD F0,0(R1)  
 LD F4,-8(R1)  
 ADD F4,F4,F2  
 SD F4,-8(R1)  
 LD F6,-16(R1)  
 ADD F6,F6,F2  
 SD F6,-16(R1)  $\rightarrow$  Hasla F16  
 ..  
 DADDI R1,R1,-64  
 BNEZ R1,loop

### 3) DAXPY loop

bop, LD F2, 0(R1)  
 MULD F4, F2, FO  
 LD F6, 0(R2)  
 ADDD F6, F4, FG  
 SD FG, 0(R2)

DADDUI R1, R1, -8  
 DADDUI R1, R1, -8  
 SLTI B3, RL, end  
 BEQZ R3, loop

ALU FP  
 ALU FP  
 LD  
 LD

ALU FP  
 SD  
 ALU FP  
 SD

3  
2  
1  
0

Fixed FO

a) Unroll 4 times

|                 |                   |                   |                   |
|-----------------|-------------------|-------------------|-------------------|
| LD F2, 0(R1)    | LD F8, 8(R1)      | LD F14, 16(R1)    | LD F20, 24(R1)    |
| MULD F4, F2, FO | MULD F10, F8, FO  | MULD F16, F14, FO | MULD F22, F20, FO |
| LD F6, 0(R2)    | LD F12, 8(R2)     | LD F18, 16(R2)    | LD F24, 24(R2)    |
| DDO F6, F4, FG  | ADD F12, F10, F12 | ADD F18, F16, F18 | ADD F24, F22, F24 |
| SD FG, 0(R2)    | SD F12, 8(R2)     | SD F18, 16(R2)    | SD F24, 24(R2)    |

DADDUI R1, R1, -32  
 DADDUI R2, R2, -32  
 SLTI R3, R1, end  
 BEQZ R3, loop

cycle 1

Mem 1  
 LD F2, 0(R1)  
 LD F14, 16(R1)  
 LD F6, 0(R2)  
 LD F18, 16(R2)

Mem 2  
 LD F8, 8(R1)  
 LD F20, 24(R1)  
 LD F12, 8(R2)  
 LD F24, 24(R2)

FP1

FP2

Int / Branch

MULD F10, F8, FO  
 MULD F22, F20, FO

DADDUI R1, R1, -32  
 DADDUI R2, R2, -32  
 SET R3, R1, end

ADD FG, F4, FG  
 ADD F12, F10, F12  
 ADD F18, F16, F18  
 ADD F24, F22, F24

SD F6, 32(R2)  
 SD F18, 16(R2)

SD F12, 8(R2)  
 SD F24, 24(R2)

Beqz R3 loop

Branch Predictors, BHR, BHTs...

1-Solo global

|    | T/NT | Predictor | HIT | BHR | Table1 | Table2 | Table3 | Table4 | Predict | HIT |
|----|------|-----------|-----|-----|--------|--------|--------|--------|---------|-----|
| B2 | NT   | 00        | Y   | 00  | 88     | 88     | 88     | 88     | NT      | Y   |
| B2 | T    | 01        | N   | 01  | 01     | 00     | 00     | 00     | NT      | N   |
| B1 | NT   | 01        | Y   | 10  | 01     | 00     | 00     | 00     | NT      | Y   |

- Las Tablas se llenan con lo que pasó en la línea nueva PERO modificando la BHT que indique el anterior

- El predictor es para el primer hit/miss

$$\begin{aligned} 00 &\rightarrow SNT \\ 01 &\rightarrow WNT \\ 10 &\rightarrow WT \\ 11 &\rightarrow ST \end{aligned}$$

(June 2014)

|    |    |       |    |    |   |    |    |  |  |            |
|----|----|-------|----|----|---|----|----|--|--|------------|
| B1 | T  | NT(m) | 01 | 11 | 1 | 01 | 10 |  |  | ← anterior |
| B2 | NT | T(m)  | 01 | 10 | 0 | 01 | 10 |  |  | → Actual   |

- Para saber si es predictor T/NT, miramos BHR anterior, y en la linea anterior buscamos la branch actual

- Nos fijamos al selector anterior a la linea actual para saber si toca global o local

- El BHR va fluyendo con entradas de ceros y unos

100% of the time

## LAB ASSIGNMENT 4

### DYNAMIC CODE SCHEDULING

#### Objectives

- Understand the behavior of the Scoreboard and Tomasulo schemes.
- Study the differences between both approaches.

#### Procedure

Execute the following MIPS program in the *Simulaz3M* simulator (version 4.12!), using the **Tomasulo's Algorithm**:

```
# data segment
    .data
array:   .float 3.0, 4.0, 9.0

# text segment
    .text
    .globl main
main:
    la $t0, array
    lwc1 $f4, 0($t0)
    lwc1 $f2, 4($t0)
    lwc1 $f0, 8($t0)
    mul.s $f6, $f0, $f0
    div.s $f0, $f6, $f2
    add.s $f0, $f4, $f4
    addi $v0, $t0, 10
    syscall
```

During the lab session, you should answer these questions:

- Provide the instruction status table when the `add.s` instruction is in the *write result* step.
- Give the same table for the Scoreboard method.
- What are the hazards solved in this example by Tomasulo, but not solved by Scoreboard?

**Note:** for the simulation runs, leave the default latencies for the functional units.

## 1. Status Table (Tomasulo) ↴

↳ → 3 primeros ciclos

|                  | Emisión | Ejecución | Escritura |
|------------------|---------|-----------|-----------|
| lwc1 f4, 0(\$t0) | 4       | 5-6       | 7         |
| lwc1 f2, 4(\$t0) | 8       | 9-10      | 11        |
| lwc1 f0, 8(\$t0) | 12      | 13-14     | 15        |
| mul.s f6, f0, f0 | 13      | 16-19     | 20        |
| div.s f0, f6, f2 | 14      | 21-27     | 28        |
| add.s f0, f4, f4 | 15      | 16-17     | 18        |

→ WAR

→ scoreboard

|                  | Emisión | Lectura | Ejecución | Escritura |
|------------------|---------|---------|-----------|-----------|
| lwc1 f4, 0(\$t0) | 4       | 5       | 6         | 7         |
| lwc1 f2, 4(\$t0) | 8       | 9       | 10        | 11        |
| lwc1 f0, 8(\$t0) | 12      | 13      | 14        | 15        |
| mul.s f6, f0, f0 | 13      | 16      | 17-20     | 21        |
| div.s f0, f6, f2 | 16      | 22      | 23-29     | 30        |
| add.s f6, f4, f4 | 22      | 20      | 24-25     | 26        |

3. Tomasulo solves WAR ↴ → mul.s f6 (solved by TOMASULO)

Information to collect

- During the lab session, you should answer these questions:
- Provide the instruction status table when the `add.s` instruction is in the *write result* step.
  - Give the same table for the Scoreboard method.
  - What are the hazards solved in this example by Tomasulo, but not solved by Scoreboard?

Not solved by scoreboard



# Unit 3

## Tomasulo

LD/SD: 1 cycle  
FPMult: 4 cycles  
FPAAdd: 3 "  
FPDiv: 6 "

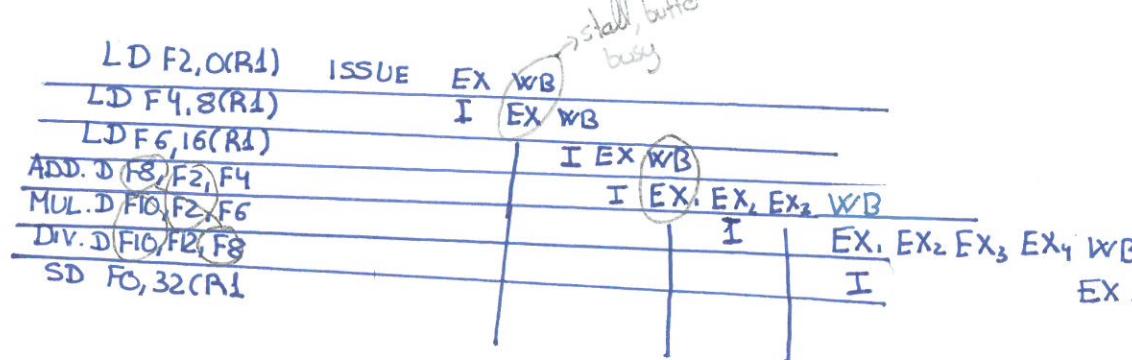
Reserv Station per FP unit: 2 entries  
LD Buffer: 2 entries  
SD Buffer: 1 entry

EX stage can't be in the same cycle in which the result is written into the CBD

### Instruc

|                  | Emission | Exec. | Escritur. |
|------------------|----------|-------|-----------|
| LD F2,0(R1)      | 1        | 2     | 3         |
| LD F4,8(R1)      | 2        | 3     | 4         |
| LD F6,16(R1)     | 4        | 5     | 6         |
| ADD.D F8,F2,F4   | 5        | 6-8   | 9         |
| MUL.D F10,F2,F6  | 7        | 9-12  | 13        |
| DIV.D F10,F12,F8 | 9        | 13-18 | 19        |
| SD F0,32(R1)     | 20       | 21    | 22        |

LD BUFF ✓ + 1 flat. + write



## Unit 4

DA DDIU R1,R0,#2

1: DADDIU R1,R0,#4

2: DADDIU R2,R2,#-1

BNEZ R2,L2

→ Run 8 times

DADDIU R1,R1,#-1

BNEZ R1,L1

→ Run 2 times

BRANCH 2  
BRANCH 1

a) 1 bit predict not taken

(\*) Takes what happens before  
1 = Taken  
0 = Not taken

|   | Times runned | Prediction | Real      | State                               |
|---|--------------|------------|-----------|-------------------------------------|
| 1 | 1            | NOT TAKEN  | TAKEN     | 1 #3                                |
| 2 | 2            | TAKEN      | TAKEN     | 1 #2                                |
| 3 | 3            | NOT TAKEN  | TAKEN     | 1 #1                                |
| 4 | 4            | TAKEN      | TAKEN     | 0 #0 → Complex condition, not taken |
| 5 | 5            | NOT TAKEN  | NOT TAKEN | 1 #3                                |
| 6 | 6            | TAKEN      | NOT TAKEN | 1 #2                                |
| 7 | 7            | NOT TAKEN  | TAKEN     | 1 #1                                |
| 8 | 8            | SAME       | TAKEN     | 0 #0 → No Toma                      |

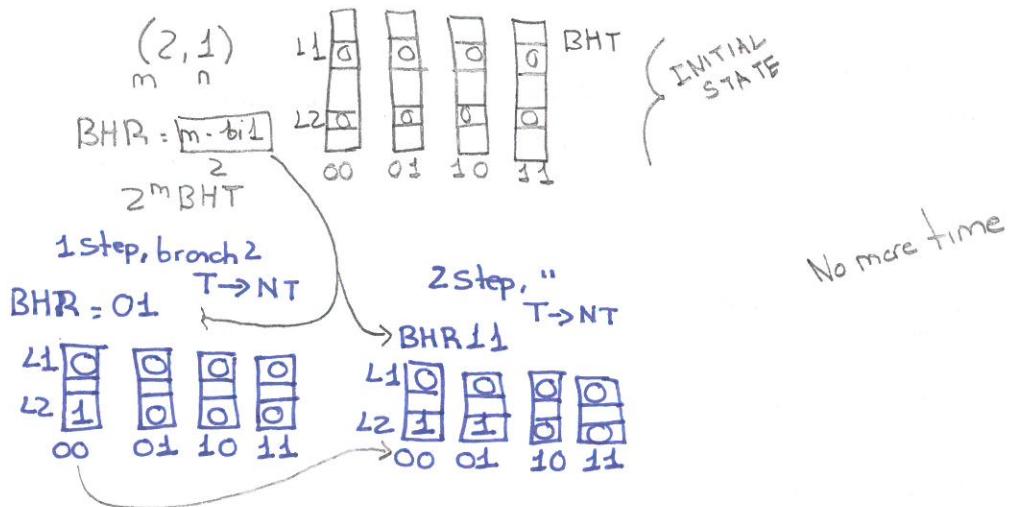
CONCLUSION: 1 BIT predict is not enough for branch 1

|   | Times runned | Prediction | Real      | STATE                    |
|---|--------------|------------|-----------|--------------------------|
| 1 | 1            | NOT TAKEN  | TAKEN     | 0111 <sup>1</sup> → 1000 |
| 2 | 2            | TAKEN      | TAKEN     | 1000 <sup>1</sup> → 1001 |
| 3 | 3            | TAKEN      | TAKEN     | 1001 <sup>1</sup> → 1010 |
| 4 | 4            | TAKEN      | NOT TAKEN | 1010 <sup>1</sup> → 1001 |
| 5 | 5            | TAKEN      | TAKEN     | 1001 <sup>1</sup> → 1010 |
| 6 | 6            | TAKEN      | TAKEN     | 1010 <sup>1</sup> → 1011 |
| 7 | 7            | TAKEN      | TAKEN     | 1011 <sup>1</sup> → 1100 |
| 8 | 8            | TAKEN      | NOT TAKEN | 1100 <sup>1</sup> → 1101 |

TAKES the prediction from the previous action, so the predict is correct

|   |           |           |      |
|---|-----------|-----------|------|
| 1 | NOT TAKEN | TAKEN     | 1000 |
| 2 | TAKEN     | NOT TAKEN | 1001 |

c) (2,1) predictor, when exec. starts, 2 last branch = NT, initial is NT



## 1.1 Computer Architecture

Comp. design task?

Maximize performance, energy efficiency...

Definition: Design to meet specific requirements, maximizing performance related to cost, power...  
Includes ISA, microarchitecture and hardware

Implementation has two components

- Microarchitecture: high-level design, memory system...
- Hardware: Detailed logic design

→ Organization: designed to meet the requirements of particular architecture. Architecture may rest, but organization evolves

Will the computer provide a multiply instruction?

Architectural

VS

Organizational  
Will the multiply instruction be implemented by a special unit or by a mechanism?

## 1.2 Performance

Latency (response time): time between start and completion of an event

Bandwidth (throughput): total work in a given time

$$\text{Performance} = \frac{1}{\text{Exec. time}}$$

"X is f% faster than Y"

$$f = (n - 1) \times 100$$

CPU exec. time = CPU clock cyc. • clock cycle time

↳ instr. count • CPI • clock cycle time

↳  $\frac{\text{CPU clock cycles}}{\text{clock freq}}$

"X is n times faster than Y"

$$\frac{\text{Perf}_X}{\text{Perf}_Y} = n$$

clock freq = 1 / clock cycle time

Instr. count CPI clock rate

|              |   |
|--------------|---|
| Program      | X |
| Compiler     | X |
| ISA          | X |
| Organization | X |
| Technology   | X |

CPI: cycles per instruc. Average num of clock cycles for each instruc.

$$\text{CPU clock cycles} = \sum_{i=1}^n (\text{CPI}_i \cdot I_i)$$

1) Comp 1: 1ns, CPI of 2 Faster?

Comp 2: 2ns, CPI of 1.2

Speed-up:  $\frac{\text{Perform. after}}{\text{Perform. before}}$

$$\frac{\text{Exec. time}_2}{\text{Exec. time}_1} = \frac{\text{Instr. num.} \cdot \text{CPI}_2 \cdot \text{clock}_2}{\text{Instr. num.} \cdot \text{CPI}_1 \cdot \text{clock}_1} = \frac{1.2 \cdot 2}{2 \cdot 1} = 1.2$$

1 is 1.2 times faster/A is 20% faster

Am-dahl's Law: Performance is improved when using some faster exec. mode is limited by the time fraction the faster mode can be used

Exec. time after enhancement = exec. time<sub>not enhan.</sub> + exec. time<sub>enhanced</sub>

2) Program runs in 100 seconds, using 80 sec. in multiplications. How much should be improve for 5 times faster?

Speed-up:  $\frac{\text{Exec. time no enhancement}}{\text{Exec. time with enhancement}} = 5$ ; New exec. time: Exec. time<sub>not enhanced</sub> + Exec. time<sub>with enhanced</sub>:  $= 20 + \frac{80}{\text{factor of enh. for mult}} \Rightarrow 5 = \frac{100}{20 + \frac{80}{\text{factor}}}$ ; IMPOSSIBLE SPEED UP OF 5

$$\text{MIPS: } \frac{\text{instruc. count}}{\text{CPU exec. time} \cdot 10^6} = \frac{\text{clock freq}}{\text{CPI} \cdot 10^6}$$

$$\text{MFLOPS: } \frac{\text{num. floating-point opera.}}{\text{CPU exec. time} \cdot 10^6}$$

Benchmark: program used to assess the performance of the whole computing system or any part

↳ Types: Real programs, kernels, toy programs

They increase predictability, set by computer vendors and periodically updated

• SPEC

• TPC

### 1.3 Classes of computers

Internet of Things (IoT), Mobile, Desktop, Servers, Clusters (Software)

Trends in AC:

- Instruction-Level Parallelism (ILP): Exploits the implicit parallelism in a code loop
- Thread-Level Parallelism (TLP): Multiple instruction dataflows that are parallel
- Data-level Parallelism (DLP): Operates on many data items at the same time
- CPU-level Parallelism
- Request-level Parallelism

F

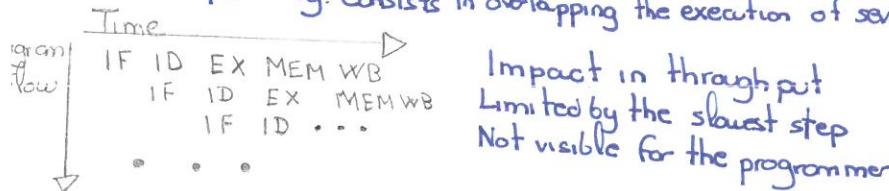
Flynn's Taxonomy: very popular but outdated. Criteria: number of instructions & data streams

→ SISD, single instruction stream, single data stream



Performance improvements:

- Pipelining: Consists in overlapping the execution of several instruc. so they use different resources



- Superscalar principles

Issue n instructions in each clock cycle    CPI = 1/n  
2-way superscalar, n=2.

IF ID EX MEM WB  
IF ID EX MEM WB  
IF ID EX MEM WB

$$CPI = 0.5$$

• VLIW (Very long instruction word): pack multiple operations in one instruction

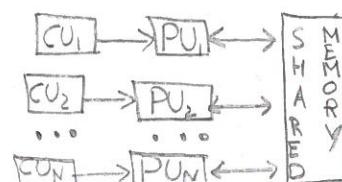
→ SIMD, single instruction stream, multiple data streams

vector architectures, multimedia extensions, graphic processor units

→ MISD, multiple instruction streams, single data stream

No commercial application

→ MIMD, multiple instruction streams, multiple data streams



Single Data Stream

SISD

MISD

Multiple Data Stream

SIMD

MIMD

Single Instruction Stream

Multiple Instruction Stream

## Increasing the ILP

15% branch, 6-7 instruc per basic block → Conclusion: Parallelism within basic blocks is limited  
↳ Solution: Exploit ILP across multiple basic blocks  
↳ How?: Overlap basic blocks and independent loop iterations

## 2.3 Study of dependences

### • Dependences:

- Property of programs
- Indicate possible hazards
- Impose an upper bound for ILP

### • Types of dependences

- Data → Antidependences
- Name
- Control → Output dependence

### Data dependences

Instructions  $i$  and  $j$  ( $i$  precedes  $j$ ),  $j$  is data dependent of  $i$  if:  
·  $i$  produces a result used by  $j$   
·  $j$  is dependent of  $k$ , and  $k$  of  $i$

Loop: LD F0,O(R1)  
ADD F4,F0,F2  
SD F4,O(R1)  
ADD R1,R1,#-8  
BNE R1,R2, Loop

### Name dependences

When two instruc use the same register or mem. location. Exists dependences if:  
·  $j$  writes a register or mem location that  $i$  reads (Antidependence)  
·  $i$  and  $j$  write the same register or mem location (Output dependence)

LD F0,O(R1)  
ADD F4,F0,F2  
SD F4,O(R1)  
LD F0,-8(R1)  
ADD F4,F0,F2  
...

### Control dependences

Ordering of an instruction  $i$  respect to a branch instruction

Every instruc except first block is control dependences

↓ ILP → Dependent instruc → possible hazards → ↑ CPI

Loop: LD F2,O(R1)  
...  
BEQ R1, exit  
LD F6,O(R1)  
ADD F8,F6,F2  
SD F8,O(R1)

### • Removing dependences

#### Name dependences

- Register renaming

Writing on registers or mem locations different than the used

Can be done statically (by the compiler) or dynamically (by the hardware)

#### Control dependences

- Removing intermediate branches
- Pending

#### Data dependences

- Scheduling → maintain the dependence, but avoid the hazard
- Transform the code

## 2.1 Pipelining basics

Pipelining: Overlap the execution of several instructions to use different resources

Two instructions can't use the same resource at the same time

Dif cache memories . Register file in ID and WB

Ideal speed-up = pipeline depth (number of steps)

Ideally: 5 stage pipeline is 5 times faster {but... sometimes next instruction cannot be executed = hazard}

Actually: steps unbalanced + overhead

Types of hazards:

- Structural hazards: Hardware cannot support a combination of instructions
- Data hazards: An instruction depends on the result of a prior one
- Control hazards: Delay

$CPI_{\text{pipelined}} = 1 + \text{pipeline stall cycles per instruc.}$  Speed-up:  $\frac{CPI_{\text{unpip}}}{1 + \text{pipeline stall...}}$

Ex: two instruc. try to write in a register with only one port

↳ Solution: stall one cycle

♦ Data hazards

Changes the order of read/write accesses to operands so that the order differs from the normal execution

Types:

RAW (read after write): Read before it has been written

WAW (write after write): Only in pipelines that write in more than one stage

WAR (write after read): Not happening in MIPS

Forwarding: Passing the AW result from who has it to who needs it

Not all data hazards can be solved by forwarding.

Some load instructions produce that problem



♦ Control hazards

Stopping the pipe: after detecting that there's a branch (after ID)

FP instructions: Multiple functional units. Some instruc. faster

|  | Branch | IF | ID | EX | MEM      | WB       |
|--|--------|----|----|----|----------|----------|
|  | +1     |    |    |    | IF stall | IF EX    |
|  | +2     |    |    |    | IF stall | IF ID EX |

|             | Latency | Initiation | • Latency: Cycles between an instruc. that produces a result and an instruc. that uses it     |
|-------------|---------|------------|-----------------------------------------------------------------------------------------------|
| Integer ALU | 0       | 1          | • Initiation interval: no of cycles that must elapse between issuing two operations of a type |
| Data mem    | 1       | 1          |                                                                                               |
| FP add      | 3       | 1          |                                                                                               |
| FP multiply | 6       | 1          |                                                                                               |
| FP divide   | 24      | 25         |                                                                                               |

## 2.2 Reducing stalls

Techniques:

- Pipelining: Overlapping instructions → The potential overlap is called instruction-level parallelism (ILP)

Goal: Reduce the CPI

How? Compilation-time analysis and Run-time analysis

Advanced (static and dynamic)

WAR, WAW, RAW

Loop unrolling → Control hazards

Branch prediction → Control hazards

Multiple instruc. per cycle → Ideal CPI

ILP related to CPI

ILP: Overlapping in a sequence of instructions, depends on the degree of dependency between instruc.

↑ ILP → few issues → few stalls → ↓ CPI

Independent instructions if they can be executed simultaneously

31) New floating point? 4 times faster than old one. Speedup of 5. Benchmark takes 100s with old one  
 ↳ time to archive that speedup?

Exec time (New) = Exec time not enhanced + Exec time enhanced

$$\text{Speedup} = \frac{\text{Exec time NE}}{\text{Exec time E}} ; 5 = \frac{100}{100-t+t/4} ; 5(100-t+t/4) = 100 ; 500 - 5t + 5t/4 =$$

c) It's not possible

$$2000 - 20t + 5t = 400$$

$$-15t = -1600$$

$$t = 106.66\dots s$$

| Type | CPI in M1 | CPI in M2 |
|------|-----------|-----------|
| A    | 1         | 2         |
| B    | 2         | 2         |
| C    | 3         | 3         |
| D    | 4         | 4         |

Evenly distributed among four types of instruc.  
 What freq for M1 to equal performance of M2 with 1000MHz  
 $\text{CPI}_{1\text{AVG}} = \frac{1+2+3+4}{4} = \frac{10}{4} = 2.5$   
 $\text{CPI}_{2\text{AVG}} = \frac{2+2+3+4}{4} = \frac{11}{4} = 2.75$

$$\frac{\text{Freq M}_1}{2.5} = \frac{1000}{2.75} ; \frac{1000 \cdot 2.5}{2.75} ; \text{Freq M}_1 = 909.0909\dots \text{MHz}$$

a) 900

Reduces CPI for FP to 35% of initial value. Improved clock cycle by 20%. Before the improvement  
 FP was 4 times the CPI for integer operations. Speed-up for a program of 6M int. instruc and 2M FP?

|         | R <sub>1</sub> | R <sub>2</sub> |
|---------|----------------|----------------|
| FP add  | 16%            | 6              |
| FP mult | 10%            | 8              |
| FP div  | 8%             | 10             |
| Non-FP  | 66%            | 3              |

b) CPU program for R<sub>1</sub> and R<sub>2</sub>. 12000 instruc

$$\text{freq} = \frac{1}{\text{tiempo}} ; \text{tiempo} = \frac{1}{\text{freq}} = \frac{1}{400 \cdot 10^6} = \dots$$

R<sub>1</sub> CPU exec = instruc count · CPI · time:

$$= 12000 \cdot 4.54 \dots = 0.0001362s$$

↳ iPhone: 2.5s - 9

R<sub>2</sub> CPU exec = instruc count · CPI · time:

$$= 12000 \cdot 13.66 \dots = 0.0004098s$$

a) MIPS for R<sub>1</sub> and R<sub>2</sub>, clock freq of 400MHz

$$R_1 = (6 \cdot 0.16) + (8 \cdot 0.1) + (10 \cdot 0.08) + (3 \cdot 0.66) = 4.54 \\ = 0.96 + 0.8 + 0.8 + 1.98 = 4.54$$

$$\text{MIPS} = \frac{\text{freq}}{\text{CPI} \cdot 10^6} = \frac{400 \cdot 10^6}{4.54 \cdot 10^6} = 88.10 \text{ MIPS}$$

$$R_2 = (20 \cdot 0.16) + (32 \cdot 0.1) + (66 \cdot 0.08) + (3 \cdot 0.66) = \\ = 3.2 + 3.2 + 5.28 + 1.98 = 13.66$$

$$\text{MIPS} = \frac{\text{freq}}{\text{CPI} \cdot 10^6} = \frac{400 \cdot 10^6}{13.66 \cdot 10^6} = 29.28 \text{ MIPS}$$

SIM CPU pg · CPI

c) What mixture of instruc R<sub>1</sub> and R<sub>2</sub> will be equally fast?

When it does not have the FP since they would have the same CPI

$$\frac{0.66 \cdot 3}{R_1} = \frac{0.66 \cdot 3}{R_2}$$

## Unit 2.

9) LB R2, 0(R1) 

DADD R3, R2, R0  
DADDUI R2, R2, S3  
SB R2, 0(R1) 4  
DADDUI R1, R1, L5  
BNE2 R3, loop 6

Data depend:  
R2 → 1 and 2  
3 and 4  
1 and 3  
R3 → 2 and 6  
R1 → none

Output depend:  
R2 → 1 and 3  
R1 → None  
R3 → None

Antidependences:  
R1 → 1 and 5  
4 and 5  
R2 → 2 and 3

P2) loop: LS F0, 0(R1)  
LS F2, 80(R1)  
MULS F4, F0, F2  
Read → SS F4, 0(R1)  
DADDUI R1, R1, #4  
BNE R1, R3, loop

P1) loop: LD R4, 0(R1)  
LD RS, 0(R2)  
SLT R4, R4, RS  
SD R4, 0(R1)  
DADDUI R1, R1, #8  
DADDUI R2, R2, #8  
BNE R1, R3, loop

## Unit 3

Ex 13/Jan 2014

Vectors x and y containing n 64 bits integer elem.  $x(i) = \begin{cases} 1 & \text{if } x(i) \leq y(i) \\ 0 & \text{if } x(i) \geq y(i) \end{cases}$

loop:

LD R4, 0(R1)  
LD RS, 0(R2)  
SLT R4, R4, RS  
SD R4, 0(R1)  
DADDUI R1, R1, 8  
DADDUI R2, R2, 8  
BNE R1, R3, loop

SLT R6, R4, RS  
SD R6, 0(R1)

start-up:

LD R7, 0(R1)  
LD R8, 0(R2)  
SLT R6, R7, R8  
LD R4, 8(R1)  
LD R5, 8(R2)

loop:

SD  
SLT  
LD  
LD  
DADDUI  
DADDUI  
BNE

finish up:

SD R6, 0(R1)  
SLT R6, R4, RS  
SD R6, 8(R1)

JAN 22

### 3.1 Static scheduling

Compiler modifies the code to increase ILP, reducing dependencies

loop unrolling → Separate dependent instruc

Software pipelining → Reorder dependent instruc

Trace scheduling → Reorder dependent instruc

#### ↳ Loop Unrolling

Static/software scheduling technique, reduce pipeline stalls due control hazards

Consists in repeating the loop code several times (reduces iterations), increase the basic block to move away dependences. In first compiling stages

|                   |                         |                         |
|-------------------|-------------------------|-------------------------|
| LD F0, 0(R1)      | LD F0, 0(R1)            | LD F0, 0(R1)            |
| ADD.D F4, F0, F2  | ADD.D F4, F0, F2        | ADD.D F4, F0, F2        |
| SD F4, 0(R1)      | SD F4, 0(R1)            | SD F4, 0(R1)            |
| LD F0, -8(R1)     | Name ADD.D F8, F6, F2   | LD F6, -8(R1)           |
| ADD.D F4, F0, F2  | Renome ADD.D F8, F6, F2 | LD F10, -16(R1)         |
| SD F4, -8(R1)     | ⇒ SD F8, -8(R1)         | LD F14, -24(R1)         |
| LD F0, -16(R1)    | ADD.D F10, -16(R1)      | ADD.D F4, F0, F2        |
| ADD.D F4, F0, F2  | ADD.D F12, F10, F2      | ADD.D F8, F6, F2        |
| SD F4, -16(R1)    | SD F12, -16(R1)         | ADD.D F12, F10, F2      |
| .D.F0, -24(R1)    | LD F14, -24(R1)         | ADD.D F16, F14, F2      |
| ADD.D F4, F0, F2  | ADD.D F16, F14, F2      | SD F4, 0(R1)            |
| SD F4, -24(R1)    | DADDUI R1, R2 #32       | SD F8, -8(R1)           |
| ADDUI R1, R1, #32 | BNE R1, R2, loop        | DADDUI R1, R1, #32      |
| 3NE R1, R2, loop  |                         | SD F12, 16(R1) → -16+32 |
| stalls ↴          |                         | BNE R1, R2, loop        |
|                   |                         | SD F16, 8(R1) → -24+32  |

7 cycles per array element

3.5 cycles per array element

#### Limitations

- Decrease in the amount of overhead amortized with each unroll
- Growth in code size (+ cache miss rate)
- Register pressure, may not be possible to allocate all the values to registers

#### ↳ Software pipelining

Scheduling performed by the compiler. Reduce stalls due to data hazards and CPI

Consists in transform a loop with dependent instruc and independent iterations into independent instruc and dependent iter

Iteration 1

Iteration 2

Iteration 3

Each iteration is composed by those 3 instructions, in this case it can be able to

|                  |                  |                  |
|------------------|------------------|------------------|
| LD F0, 0(R1)     | LD F0, -8(R1)    | LD F0, -16(R1) → |
| ADD.D F4, F0, F2 | ADD.D F4, F0, F2 | ADD.D F4, F0, F2 |
| SD F4, 0(R1)     | SD F4, -8(R1)    | SD F4, -16(R1)   |



## Práctica 5

Cargamos el código

```

ADDI R2, R0, 10 //Dir x
ADDI R3, R0, 30 //Dir y
ADDI R4, R0, 5 //Dir z
LF F0, (R4)
ADDI RS, R2, 10
Buc: LF F1 (R2)
LF F2 (R3)
MultF F3, F1, F2
ADDF F1, F1, F0
SF F1, (R3)
ADDI R2, R2, 1
ADDI R3, R3, 1
BNE R2, R5, bnc
SW RS, 0(R0)
SW R4, 1(R0)
    { 1st loop

```

Load mem

- #GPR
- #FPR
- #MEM
- [S1] 2
- [10] 1 2 3 4 5 6 7 8 9 10
- [31] 1
- [33] 1
- [35] 1
- [37] 1
- [39] 1

- a) Behaviour of branch predict. Used by the simulator. State in every branch instruc
- Uses dynamic predictor of two bits. Start at 00 then 01 | 10 00.
  - First loop fails so  $00 \leftarrow 1; 01$  ↳ 00 means (predicts not taken, but is taken 10 times)
  - Second loop is not taken, but should have been taken, so  $01 \leftarrow 1; 11$
  - Third, far... loops are not taken, all 01 until last, that is not taken so  $11 \leftarrow 0; 10$

00  
01  
11  
10

b) Predictor hit rate

$$\text{Fails: } 1^{\text{st}}, 2^{\text{nd}}, \text{last} \quad \text{Hit}_R = \frac{7 \text{ hit}}{11 \text{ instruc}} = 63.60$$

Total:  $1 + 10 = 11$

c) Prefetch, decoder and reorder in cycles 26 and 27

26 Prefetch

Decoder  
mmmm  
- ? ?  
? ?

| Entry | Instruc | Dest. | Value | Address | State |
|-------|---------|-------|-------|---------|-------|
| 11    | 11      | 3     | 31    | -1      | write |
| 12    | 12      | -1    | 1     | -1      | write |
| 13    | 13      | -1    | 20    | 0       | write |
| 14    | 14      | -1    | 5     | 1       | write |

No commit  
bc fails predict

Fail of first predict. Prefetch and decoder empty bcs when fails it doesn't keep exec. next instruc of next loop until first ends. Empty bcs only left last 4 writing stages of last iterat. (Nothing in deco or prefetch)

27 Prefetch

Decoder

Entry Instruc Dest. Value Address State

Vacio de

5

6

7

~ 8

System starting next iteration for second loop. It will end as first loop in cycle 26. Some in last



# First Part, Superscalar

a) CPI for degree of 2

$$CPI = \frac{141 \text{ cycles}}{6 + (8 \cdot 16)} = \frac{141}{134} = 1.0597$$

b) CPI for degrees of 4, 6, 8, 16

$$CPI_4 = \frac{137}{6 + (8 \cdot 16)} = \frac{137}{134} = 1.02985$$

$$CPI_6 = \frac{134}{6 + (8 \cdot 16)} = \frac{134}{134} = 1$$

$$CPI_8 = CPI_{16} = CPI_6$$

c) What aspect that is limiting performance? Improvement to best speedup?

Seen before, as higher degree, a better CPI (closer to 0)

Degree 2 = 0.5; Degree 4 = 0.125... ...

Cannot happen due to program dependencies

Best option is degree 6, since we obtain the ideal CPI of 1.

⇒ d)

GWR

$61282 \rightarrow R3$

$2S12$

$36152$

$100$

$1SF$



$226$

$307$

$402$

$308$

$0$

$6$

$7$

$152$

$10$

$8$

$12$

$9 \quad 13$

$11$

|              | Num | Lat. |
|--------------|-----|------|
| Suma         | 1   | 2    |
| Mult         | 2   | 3    |
| Suma F       | 1   | 3    |
| Mult F       | 2   | 4    |
| Memoria      | 2   | 2    |
| Fallo cache: |     | 9    |
| Salto:       | 1   | 2    |

Second part, VLW

↳ Debe de dar 13

|   |   |     |
|---|---|-----|
| 3 | 0 | 226 |
| 4 | 2 | 307 |
| 5 | 1 | 402 |
| 6 |   | 308 |

|    |   |    |
|----|---|----|
| 7  |   |    |
| 10 | 8 |    |
| 12 |   |    |
| 11 | 9 | 13 |

Bucle:

$LF \quad F6, (R2)$   
 $MULTF \quad F6, F6, F6$   
 $MULTF \quad F6, F6, F1$   
 $SF \quad F6, (R12)$   
 $ADDI \quad R2, R2, \#1$   
 $ADDI \quad R12, R12, \#1$   
 $ADDI \quad R3, R3, \#1$   
 $BNE \quad R3, R1, Bucle$

Lectura elem. v  
 elem al cuadrado  
 2. m. elem al cuadrado

d)

LF F6, 0(R2)  
LF F1, 1(R2)  
LF F4, 2(R2)  
LF FS, 3(R2)

MULTF F6, F6, F6  
MULTF F1, F1, F1  
MULTF F4, F4, F4  
MULTF FS, FS, FS  
ADDI R3, R3, #4

MULTF F6, F6, F8  
MULTFF F1, F1, F2  
MULTFF F4, F4, F3  
MUTF FS, FS, F1  
SF F6, -4(R12)  
SF F1, -3(R12)  
SF F4, -2(R12)  
SF FS, -1(R12)  
ADDI R2, R2, #4  
ADDI R12, R12, #4  
SNE R3, R12, bckle

$$CPI = \frac{98 \text{ cycles}}{6 + (4 \cdot 20)} \cdot 1.1$$

↙ Loops      ↘ instruc

Initially: 16 loops  
After unrolling: 4 loops }  $16/4 = 4$

Instruc BF loop init: 6

Instruc AF loop init: 6

Instruc per loop with NO LU: 8  
Instruc per loop with LU: 20

Total before LU: 134

Total after LU: 86

| Instruction      | Issue | Execute | Write Result | Commit |
|------------------|-------|---------|--------------|--------|
| 0 ADDI R2 R0 #10 | 3     | 4       | 5            | 6   4  |
| 1 ADDI R3 R0 #30 | 3     | 5       | 6            | 7   4  |
| 2 ADDI R4 R0 #5  | 5     | 6       | 7            | 8   5  |
| 3 LF F0 (R4)     | 7     | 7       | -            | 8      |
| 4 ADDI R5 R2 #10 | 7     | 8       | 8            | 9   6  |
| 5 LF F1 (R2)     | 8     | 8       | 8            | 9   6  |
| 6 LF F2 (R3)     | 8     | 8       | 8            | 10   6 |
| 7 MULT F1 F1 F2  | 7     | 7       | 7            | 11   6 |
| 8 ADDF F1 F1 F0  | 8     | 8       | 8            | 11   6 |
| 9 SF F1 (R3)     | 8     | 8       | 8            | 11   6 |
| 10 ADDI R2 R2 #1 | 9     | 9       | 9            | 12   9 |
| 11 ADDI R3 R3 #1 | 9     | 9       | 9            | 12   9 |
| BNE R2 R5        | -     | -       | -            | -      |
| SW R5 0(R0)      | -     | -       | -            | -      |
| SW R4 1(R0)      | -     | -       | -            | -      |

F0  
R4 = 5

|    |    |    |    |    |
|----|----|----|----|----|
| 0  | 3  | 4  | 5  | 6  |
| 1  | 3  | 5  | 6  | 7  |
| 2  | 5  | 6  | 7  | 8  |
| 3  | 5  | 7  | 13 | 14 |
| 4  | 6  | 7  | 8  | 14 |
| 5  | 6  | 7  | 13 | 15 |
| 6  | 7  | 8  | 14 | 15 |
| 7  | 7  | 14 | 14 | 15 |
| 8  | 8  | 20 | 24 | 21 |
| 9  | 8  | 9  | 25 | 26 |
| 10 | 9  | 10 | 11 | 26 |
| 11 | 9  | 11 | 12 | 27 |
| 12 | 10 | 11 | 13 | 27 |
| 13 | 10 | 11 | 13 | 27 |
| 14 | 11 | 12 | 14 | 27 |

